

Paper Title

1st Given Name Surname

dept. name of organization (of Aff.) name of organization (of Aff.) City, Country email address or ORCID

2nd Given Name Surname

dept. name of organization (of Aff.) name of organization (of Aff.) City, Country email address or ORCID

Abstract—

- **Problem:** improve settling prediction for a position-dependent dual gantry without surrendering the meaning of its physical parameters
- **Method:** combine a quasi-LPV first-principles baseline, dynamic parallel LFR augmentation, closed-loop rollout training, and state-level orthogonal projection
- **Evidence:** compare baseline, unconstrained augmentation, orthogonal variants, and a black-box model on the completed validation ladder, including relevant motion and settling where available
- **Main result:** report predictive improvement together with the measured limits from short-horizon training, inactive added states, and non-orthogonal true missing dynamics
- **Conclusion:** state when the proposed decomposition is useful and which conditions prevent interpreting it as true parameter recovery

[TODO: Write the numerical abstract only after the final simulation and Telica result tables are frozen.]

Index Terms—grey-box identification, LPV systems, model augmentation, orthogonal projection, dual gantry

I. INTRODUCTION

[TODO: Kijken bronnen origineel van roland + hoekstra + bendeguz + roel drenth. currently i dont really state anything about orthogonality for interpretability for example.]

A. Industrial motivation and modelling gap

- Settling prediction limits throughput in coupled high-precision gantries, so model error has a direct operational cost [1]
- The concrete digital-twin use case is to evaluate relevant ASMPT motion and a changed controller before deployment on the machine
- The physics model exposes masses, stiffnesses, damping, and coupling, but omits position-dependent flexible dynamics that matter for prediction
- A black-box model offers flexibility but weakens extrapolation arguments and removes parameters useful for control design, digital twins, and health monitoring
- Grey-box dynamic augmentation is the target compromise: preserve the baseline and learn only the missing state dynamics [2]

[TODO: Which industrial benefit should lead: settling prediction, controller design, offline commissioning, or condition monitoring?] [TODO: Instead of health monitoring more focus

on process and root cause analysis, especially the interpretability is waardevol for this.] [TODO: Clearly explain white, black and grey-box] [TODO: Clearly explain interpretability and should be clear in orthogonality section also. lets see where to say what..] [TODO: Decided 2026-10-06: the general definition of interpretability belongs here; Section III states only its operational form (the estimated θ_{base} keeps its physical meaning, which requires a unique estimate) and does not repeat a general one. Candidate sources: Györök et al. [3], Sec. 1 and 3 (“With non-unique θ_b , the baseline model could lose its physical meaning”); Kessels [4], Sec. 5.3.1.3 (negation) and Sec. 6.2 (orthogonality “could enhance physical interpretability”, not yet available beyond linear-in-the-parameter problems), which also states the gap.] ASMPT develops semiconductor assembly and packaging equipment, including high-precision motion systems for advanced packaging. Advanced packaging refers to combining multiple semiconductor functions within a single package, for example through heterogeneous integration of different die types such as logic and memory, and through hybrid bonding, in which dies are directly joined at very fine pitch. [TODO: Shorten to a clause inside the previous sentence (page budget).] Such technologies place strict demands on alignment accuracy, yield, throughput, and manufacturing cost [5], [6]. [TODO: Only partly supported. IRDS covers heterogeneous integration, cost and yield, but never mentions hybrid bonding and its throughput means memory bandwidth. Workman covers sub-micron alignment and yield. Neither supports machine throughput: drop it here or add a source. Workman bib entry: authors incomplete, pages 2071 to 2077 missing.]

This project focuses on a dual-gantry motion system [TODO: Tense and terminology: project becomes paper.] in which two parallel linear actuators support a common beam carrying the payload. The shared beam couples the actuator dynamics, so motion of one actuator influences the other, and the system exhibits position-dependent behaviour and flexible dynamics [1], [7]. [TODO: Teo (abstract only) supports coupling and payload-position dependence, not flexibility. Flexibility rests on Garcia Sec. 2 (flexible joints kb1, kb2): consider flexible joint dynamics.] In such systems, settling behaviour is a key performance metric. After each motion, the system must wait for residual oscillations to decay before the next placement can begin [8], so settling time directly affects throughput. Furthermore, cross-coupling means that the motion of one gantry can disturb the other, making accurate prediction relevant

not only for settling but also for tracking behaviour during simultaneous motion.

Improving settling performance therefore requires an accurate representation of the underlying plant dynamics. ASMPT currently captures part of the position dependence of the dual-gantry system using a baseline model based on frequency-response measurements at multiple operating points with interpolation across the operating space. [TODO: No written source: confirm with Quinten or Jasper. The paper does not compare against this model, so do not imply it.] However, this baseline captures only part of the position dependence and does not systematically account for cross-coupling between the gantries. [TODO: Same confirmation as above.] A more accurate plant model is therefore needed. Such models are not only valuable for settling prediction, but also for digital-twin development, offline evaluation, control design, and condition monitoring. [TODO: Lead with the controller-evaluation digital twin: test a changed controller on the model before the machine (Quinten, 18 Sep). Condition monitoring is not evaluated; its supporting sentence was removed.] Digital twins can support offline simulation of semiconductor manufacturing equipment and help optimize operating parameters before testing changes on the real machine, reducing costly trial-and-error [9]. [TODO: Only partly supported (abstract only). Hwang optimises operating parameters of wafer-processing cluster tools by digital-twin simulation; before testing on the real machine and trial-and-error are not in the source.]

White-box models are attractive because they reflect the physics of the system, which makes them interpretable. First-principles (FP) models are a specific type of white-box model, derived from physical laws and described by quantities such as mass, stiffness, and damping. However, purely physics-based models do not always capture all behaviour needed for accurate prediction across the operating range [10]. For the dual-gantry system, this limitation is also evident in the first-principles model structure of García-Herreros et al. [1], which is only able to capture part of the position-dependent behaviour and cross-coupling. [TODO: Not what Garcia says. Garcia Sec. 2 deliberately neglects cross-arm vibration on the flexible plates, detent forces, and friction variation along the stroke. Our LPV-LFR baseline captures the $M(Y)$ position dependence exactly; the missing dynamics here are the hidden absorber mode and, in the friction twin, Coulomb friction in the truth only (D-204).] [TODO: Proposed replacement of the Garcia sentence: “For the dual-gantry system, the first-principles model of García-Herreros et al. neglects the resonance of the supporting base, the vibration of the cross-arm on its flexible plates, the detent forces of the linear motors, and the variation of friction along the stroke [1, Sec. 2.4].” Verified against Garcia Sec. 2.4, pp. 15 to 16. Our baseline additionally omits Garcia’s Coulomb friction (stated in Sec. II); the simulated stand-ins (absorber, friction) are introduced in Setup.]

Purely data-driven models such as neural networks or Gaussian processes can describe measured behaviour well, but they give less physical insight, are often harder to trust [TODO: Grammar: an and is missing after insight.] at operating

conditions not covered by the training data [10]. Typically these models also require large amounts of data and significant computational power to train. [TODO: Uncited. Large data requirements: Willard Sec. 1. Computational power: no source found, drop or cite. Verified alternative: black-box models spend effort rediscovering dynamics already known from first principles (Hoekstra 2025 EJC p. 1).] For the plant model of the dual-gantry system, both predictive accuracy and physical interpretability are required. However, neither a purely physics-based nor a purely data-driven model is sufficient on its own [10], [11]. [TODO: Cite Willard only: Noel never argues that pure physics models are insufficient. Willard bib: ACM Computing Surveys 2022 (issue 2023), not 2020; month field wrong.] Grey-box modeling is the logical choice, because it is able to combine a physically meaningful baseline model needed for interpretation with data-driven flexibility needed for accurate prediction [11]. [TODO: Noel supports only grey-box combining physical structure with data-driven flexibility. Name model augmentation explicitly and cite Hoekstra 2025 EJC and Hoekstra 2026.] However, related work at ASMPT shows that preserving physical interpretability during model updating remains an open challenge [4]. [TODO: Supported by Kessels thesis Ch. 6 recommendations (PDF p. 222): augmentation can compensate for incorrect FP parameters, compromising their interpretability, and orthogonality methods exist only for linear-in-the-parameter cases. The thesis says industrial wire bonder and never names ASMPT: confirm or reword.]

Learning-based hybrid modeling has already been demonstrated on real industrial systems [12]. [TODO: Supported (Retzler 2024, industrial robot, measured data).] However, a key gap remains: there is no validated grey-box modeling framework for the dual-gantry system that simultaneously addresses coupled multi-input multi-output (MIMO) dynamics, position-dependent behaviour, and preserves physical interpretability. [TODO: Rewrite the gap as a verified intersection. Hoekstra 2025/2026: dynamic LFR augmentation, no LPV scheduling, no closed-loop training. Kessels: closed-loop-trained augmentation on an industrial motion system, but no LFR or LPV, frozen FP parameters, no orthogonality, so closed-loop training alone is not the novelty. Gyrook OBC 2026: discrete-time input-output models only; its conclusion names state-space without full-state measurement as future work. Validated must be scoped to simulation.] This project therefore aims to develop and validate such a framework [TODO: Tense and scope: this paper develops and evaluates in simulation.] so that the plant model becomes more accurate for settling prediction while remaining physically meaningful.

B. Research question and scope

- Main question: how should a physics-based plant model be augmented to capture unmodelled dynamics, improving prediction accuracy across operating points and motion profiles while preserving the physical interpretability of its parameters?
- Aspect 1 asks in what form the baseline is represented and whether the position-dependent (LPV) form predicts better

than position-independent (LTI) alternatives; answered by the derivation plus a small demonstration

- Aspect 2 asks which augmentation structure effectively captures the unmodelled dynamics: parallel with or without added states (parallel versus series argued in Section III), and what makes the added states learn them
- Aspect 3 asks how parameter interpretability is maintained compared to unconstrained joint estimation; answered by what orthogonality by construction guarantees and what it cannot
- Aspect 4 asks how well the model generalises to held-out operating points and unseen motion profiles compared to a black-box state-space model; controller transfer is shown as a use case
- Scope: simulation only; aspects are answered by derivation where possible and by experiment where needed

[**TODO:** Outline bullets updated to the changed research question and aspects (D-241).]

The main research question is formulated as follows:

For the dual-gantry system, how should a physics-based plant model be augmented to capture unmodelled dynamics, improving prediction accuracy across operating points and motion profiles while preserving the physical interpretability of its parameters?

[**TODO:** changed rq and polished aspects] To answer this research question, the following four aspects are investigated:

1) Position-dependent baseline modeling

In what form should the baseline model be represented for subsequent augmentation, and does position-dependent (LPV) formulation improve prediction compared to position-independent (LTI) alternatives? [**TODO:** LPV versus frozen LTI exists only as a baseline check (lpv_lfr_baseline/scripts/validate_lfr.py), not in the final evaluation; Quinten wants LTI as second baseline.]

2) Augmentation structure

What augmentation structure can effectively capture dynamics not modeled in the physics-based baseline model?

3) Interpretability preservation

How can parameter interpretability be maintained during data-driven augmentation, compared to unconstrained approaches?

4) Model generalization

How well does the augmented model generalize to (i) held-out operating points and (ii) unseen motion profiles, and how does its prediction accuracy compare to a black-box state-space model? [**TODO:** Not yet true. The D-196 black-box arm runs on the MA_FRAC 0.10 dataset, not the production ma50 data. Unseen motion profiles partly covered by the Telica cycloidal records (D-206). Controller transfer is not implemented. Align with the validation ladder in Writing/README.md.]

C. Contributions and paper organisation

- Derive a compact continuous-time LPV-LFR realisation of the dual-gantry baseline with rational dependence on payload position
- Extend dynamic parallel LFR augmentation to a closed-loop, self-scheduled state-space setting with learned additional states
- Construct state-level orthogonality for the identifiable physical-parameter tangent space and quantify its true-parameter recovery condition
- Evaluate predictive fit, settling, controller transfer, added-state contribution, interpretability, and black-box performance under matched data
- Paper organisation: baseline, augmentation, OBC, setup, results, discussion, conclusion, then supporting derivations
- Figure plan: no dedicated introduction figure; use the limited figure budget where the baseline, training interconnection, and OBC need visual explanation

[**TODO:** Confirm whether references and appendices are excluded from the 12-page IEEE limit and whether TU/e requires front matter beyond the IEEE paper.]

II. DUAL-DRIVE GANTRY SYSTEM AND PHYSICS BASELINE

A. System and first-principles baseline

The physics baseline of the gantry system used in this paper is the lumped-parameter model of the the dual-drive gantry of García-Herreros et al. [1]. The gantry has two actuators that are connected with a cross-arm and a third actuator that moves the payload along this arm (Fig. 1). [**TODO:** Does this make sense to put here?: Dynamics not represented by the baseline are treated as model discrepancy and addressed by the augmentation in Section III.] Define the actuator forces $u_{\text{act}} = [F_{X1} \ F_{X2} \ F_Y]^\top$, the measured positions $q_{\text{act}} = [X_1 \ X_2 \ Y]^\top$, and the generalized coordinates $q = [X \ \Theta \ Y]^\top$. Let L_b denote the cross-arm length between




Fig. 1: [**TODO:** Lumped-parameter gantry schematic adapted from the mechanical arrangement in [1]. Red labels denote actuator coordinates and blue labels generalized coordinates.]

the two flexible joints. The coordinate and force transformations are

$$q = P^{-\top} q_{\text{act}}, \quad u = P u_{\text{act}}, \quad (1)$$

$$P = \begin{bmatrix} 1 & 1 & 0 \\ L_b/2 & -L_b/2 & 0 \\ 0 & 0 & 1 \end{bmatrix},$$

where u denotes the generalized force. The generalized coordinates are defined as $X = (X_1 + X_2)/2$ and $\Theta = (X_1 - X_2)/L_b$, while Y is unchanged. These coordinates avoid [TODO: This needs to make clear why we choose for this realization. a closed kinematic-chain description and make the coupling caused by the nonuniform load distribution explicit] [1]. [TODO: How to connect, nonuniform ... moving payload position] Moving the payload position Y changes the mass distribution of the gantry and therefore its inertia. This position dependence is retained by representing the baseline in linear parameter-varying (LPV) form, with the measured payload position Y as scheduling variable [13].

Following García-Herreros et al. [1], the baseline uses $\cos \Theta \approx 1$ and $\sin \Theta \approx 0$ and assumes that the Coriolis and centripetal force term is negligible compared with the inertia, damping, and stiffness forces. Under these assumptions, the baseline equations of motion are

$$M(Y)\ddot{q}(t) + C\dot{q}(t) + Kq(t) = Pu_{\text{act}}(t), \quad (2)$$

with

$$M(Y) = \begin{bmatrix} m_1 + m_2 + m_b + m_h & \frac{L_b}{2}(m_1 - m_2) - m_h Y & 0 \\ \frac{L_b}{2}(m_1 - m_2) - m_h Y & J_b + J_h + \frac{L_b^2}{4}(m_1 + m_2) + m_h(d^2 + Y^2) & -m_h d \\ 0 & -m_h d & m_h \end{bmatrix}, \quad (3)$$

$$C = \begin{bmatrix} c_{g1} + c_{g2} & \frac{L_b}{2}(c_{g1} - c_{g2}) & 0 \\ \frac{L_b}{2}(c_{g1} - c_{g2}) & c_{b1} + c_{b2} + \frac{L_b^2}{4}(c_{g1} + c_{g2}) & 0 \\ 0 & 0 & c_y \end{bmatrix}, \quad (4)$$

$$K = \begin{bmatrix} 0 & 0 & 0 \\ 0 & k_{b1} + k_{b2} & 0 \\ 0 & 0 & 0 \end{bmatrix}.$$

[TODO: Proposed insertion: “García-Herreros et al. also model Coulomb friction at the three actuators [1, Sec. 2.2]. The baseline omits these terms, so that friction remains a physical effect the augmentation must capture (Section V).” Reason documented in D-204. Second reason? A sign term is not twice differentiable, which [2, Cond. 8] assumes for well-posedness. because: ?] [TODO: Look at this section, the above state a lot, look fine] The masses, flexible-joint stiffness and damping coefficients, and [TODO: geometric dimensions] are shown in Fig. 1. Here, J_b and J_h denote the yaw inertias, while c_{g1} , c_{g2} , and c_y denote the actuator damping coefficients. The nominal values are taken from [1, Table I] and listed in Appendix A-B [TODO: Not sure if to add this, depends on real data or not; they were not identified on the present Telica machine.] To prevent the definition of q from being parameter-dependent, the cross-arm length L_b is fixed at the nominal value listed in Appendix A-B rather than treated as a free model parameter. [TODO: Determine and justify the operating




Fig. 2: [TODO: Self-scheduled baseline interconnection. The Delta-block closes the latent variables of G through $w = \Delta(Y)z$, with $\Delta(Y) = YI_6$ and Y a component of the state.]

range of Y ; distinguish the mechanical stroke, simulated-data range, and range covered by the measured ILC records.]

B. Self-scheduled LPV-LFR representation

- Y is a state and the scheduling variable; $M(Y)^{-1}$ makes the model rational in Y (Drenth)
- $M(Y) = M_0 + Y M_1 + Y^2 M_2$; Y enters only through $Y\ddot{q}$ and $Y^2\ddot{q}$, realized by repeating Y (this work)
- Latent variables z, w and Delta-block $\Delta(Y) = YI_6$; not claimed minimal (this work)
- LTI part G in the continuous-time form of Drenth Eq. (6); its first block row is the state equation (adapted)
- Closing the algebraic loop gives back (2) exactly (this work)
- Well posed iff $M(Y)$ nonsingular, true for every Y ; closed-form $M(Y)^{-1}$, no numerical solve (Drenth Def. 1; this work)

The measured payload position Y is both a state and the scheduling variable, making the baseline a self-scheduled LPV model [TODO: source, or do i state it earlier..?]. Solving (2) for $\ddot{q}$ requires $M(Y)^{-1}$, so the model is rationally dependent on Y . An Linear Fractional Representation (LFR) is introduced to separate this rational scheduling dependence. We move the dependence out of the LTI part G and into a Delta-block $\Delta(Y)$, with the scheduling being connected in feedback with G (Fig. 2) [2], [14].

To separate the scheduling variable Y , deconstruct the inertia matrix (3) into its Y dependent components,

$$M(Y) = M_0 + Y M_1 + Y^2 M_2. \quad (5)$$

M_0 , M_1 and M_2 are given in Appendix A-A. Substituting (5) into (2) gives

$$M_0\ddot{q} + M_1(Y\ddot{q}) + M_2(Y^2\ddot{q}) = u - C\dot{q} - Kq. \quad (6)$$

Here Y only enters through $Y\ddot{q}$ and $Y^2\ddot{q}$, with $Y^2\ddot{q} = Y(Y\ddot{q})$. This lets us apply the same scheduling variable twice instead of introducing a second scheduling variable Y^2 ,

$$\ddot{q} \xrightarrow{Y} a_1 \xrightarrow{Y} a_2, \quad (7)$$

with $a_1 = Y\ddot{q}$ and $a_2 = Y^2\ddot{q}$.

[TODO: Left here, ...] We treat a_1 and a_2 as independent signals (6) becomes

$$M_0\ddot{q} + M_1a_1 + M_2a_2 = u - C\dot{q} - Kq, \quad (8)$$

which contains no Y . The **latent variables** are $z = \text{col}(\ddot{q}, a_1)$ and $w = \text{col}(a_1, a_2)$, with the **Delta-block** mapping z to w .

$$w = \begin{bmatrix} a_1 \\ a_2 \end{bmatrix} = \begin{bmatrix} Y\ddot{q} \\ Ya_1 \end{bmatrix} = \underbrace{\begin{bmatrix} YI_3 & 0 \\ 0 & YI_3 \end{bmatrix}}_{\Delta(Y)=YI_6} \begin{bmatrix} \ddot{q} \\ a_1 \end{bmatrix} = \Delta(Y)z. \quad (9)$$

We do not claim this **realization** to be **minimal** **[TODO: I have an SVD reduction add this in the appendix or not?]**. We can write the LTI part G of the LFR in the continuous-time form, [14, Eq. (6)],

$$\begin{bmatrix} \dot{\tilde{x}} \\ z \\ q \end{bmatrix} = G \begin{bmatrix} \tilde{x} \\ w \\ u \end{bmatrix}, \quad (10)$$

where $\tilde{x} = \text{col}(q, \dot{q})$ is the **state** and q the **output**. The first block row is the **state equation**, with $\ddot{q}$ obtained from (8) by multiplying with M_0^{-1} . For the blocks of G Appendix A-A. Closing the **algebraic loop** with (9) results in $a_1 = Y\ddot{q}$ and $a_2 = Y^2\ddot{q}$, which turns (8) back into (2). The LFR is therefore an exact realization of the baseline. The LFR is **well posed** if the **algebraic loop** has a unique solution for every Y [14, Def. 1]. For the gantry realization we show that this conditions holds when $M(Y)$ is nonsingular (Appendix A-A). For positive **masses** and **inertias**, $M(Y) \succ 0$ for every $Y \in \mathbb{R}$ (Appendix A-C). Therefore the **LFR** is well posed over the whole Y -operating range. **[TODO: Connect the two sentences better?]**

C. Physical parameters and identifiable combinations

- The input-output behaviour fixes M_0, M_1, M_2, C, K and conversely, so the identifiable part of θ is what the matrices contain (this work)
- Twelve nonzero entries, three equal to m_h , leave ten independent functions: exactly ten combinations are structurally identifiable (this work; definition Hong, Ovchinnikov)
- Written physically as θ_{base} , the gantry counterpart of base parameters (Gautier and Khalil); four lost directions; estimates reported in θ_{base} ; recovery from the records in Setup
- Positive θ_{base} does not give $M(Y) \succ 0$; that holds iff (12) (this work; proof Appendix A-C)
- Training keeps (12) at every iterate; the map and its justification are in Section III-D

The fourteen raw parameters θ of Table I cannot all be identified from input-output data. Two parameter vectors with the same input-output behaviour satisfy

$$(M(Y) - M'(Y))\ddot{q} + (C - C')\dot{q} + (K - K')q = 0$$

along every trajectory. Since every $(q, \dot{q}, \ddot{q})$ is attained, the input-output behaviour fixes M_0, M_1, M_2, C, K , and conversely. The identifiable part of θ is therefore what these matrices contain.

Up to symmetry, (3) and (4) have twelve nonzero entries. Three of them, $M_{0,33}$, $-M_{1,12}$ and $M_{2,22}$, equal m_h , which leaves ten distinct functions of θ . Their Jacobian with respect to $(m_h, d, m_1, m_b, J_b, c_{g1}, c_{g2}, c_{b1}, c_y, k_{b1})$ is block triangular with determinant $m_h L_b^2/2 \neq 0$, so the ten are independent. Exactly ten combinations of θ are therefore structurally identifiable [15, Def. 2.5], [16, Def. 7].

We write them in physically readable form as

$$\theta_{\text{base}} = [k_{b,\Sigma}, c_{g1}, c_{g2}, c_y, c_{b,\Sigma}, m_h, m_\Sigma, m_\Delta, J_{\text{eff}}, d]^\top, \quad (11)$$

with $k_{b,\Sigma} = k_{b1} + k_{b2}$, $c_{b,\Sigma} = c_{b1} + c_{b2}$, $m_\Sigma = m_1 + m_2 + m_b$, $m_\Delta = m_1 - m_2$ and $J_{\text{eff}} = J_b + J_h + \frac{L_b^2}{4}(m_1 + m_2)$, the gantry counterpart of base inertial parameters [17]. The four lost directions are $k_{b1} - k_{b2}$, $c_{b1} - c_{b2}$, $J_b - J_h$ and a mass transfer from the carriages to the cross-arm, so estimates are reported in θ_{base} . Recovery from the closed-loop records is examined in Section V. Positive θ_{base} does not imply $M(Y) \succ 0$. For $m_h, m_\Sigma, J_{\text{eff}} > 0$, $M(Y) \succ 0$ for every Y if and only if

$$m_\Sigma J_{\text{eff}} > \frac{L_b^2}{4} m_\Delta^2 \quad (12)$$

(Appendix A-C). Training therefore uses logarithmic coordinates for the nine positive combinations and a tanh-bounded m_Δ (Section III-D), so (12) and well-posedness hold at every iterate.

III. DYNAMIC AUGMENTATION OF THE LPV-LFR BASELINE

- The section extends the baseline with learned dynamics and identifies the result from closed-loop records; it names its five subsections (style profile)
- Contribution share: the realisation on the self-scheduled gantry baseline with joint estimation of θ_{base} , its identification in closed loop with the known controller, and the PS initialisation of the added states (MUST ESTABLISH)

The dynamics of the system that the baseline defined in Section II-B does not capture, will be identified from input-output data measured in closed-loop **[TODO: To vague?]**. The baseline model will be augmented to represent these missing dynamics. To represent **[TODO: unmodeled dynamics]** a correction to the states of the baseline models does not suffice **[TODO: why?]**, additional states are therefore introduced.

A. Augmented model

- If the omitted effects are dynamic, the six baseline states cannot represent them; refitting θ_{base} or a static augmentation adds no states (EJC Sec. 2)
- Hoekstra's LFR augmentation structure with a fixed interconnection; the self-scheduled LPV-LFR is its baseline block; the learned block is a network, not Drenth's LPV-LFR augmentation (thesis Eq. (5.2)) (? : reason)
- Parallel because the orthogonal-by-construction method of Section IV is defined for the additive structure (Györök 2026); output map not augmented (Kessels Remark 5.3)
- The structure is a discrete-time state update (EJC Sec. 2, Table 1), so f_{base} is one RK4 step with the input held and $M(Y)$ per stage; the correction is added after the step

- Acyclic interconnection plus the baseline result of Section II: well posed for every Y (EJC Thm. 1)
- Normalisation: affine, statistics of the training records, velocities from differenced positions; later displays in normalised symbols
- One tanh MLP supplies f_{aug} and g_{aug} (EJC Sec. 4.3); it corrects every physical row, so a persistent correction on X or Y accumulates in position (? : justification); $n_{\bar{x}} = 0$ is the static parallel augmentation
- The baseline/network split is not unique (2026 Sec. 5.2); interpretable means the estimated θ_{base} keeps its physical meaning (Györök Sec. 3); Section IV makes the estimate unique

If the effects that the baseline omits are dynamic, its six states cannot represent them. Refitting θ_{base} adds no states. Neither does a static augmentation, which corrects the state update of the baseline without extending its state [18, Sec. 2]. The augmentation therefore needs states of its own.

We use the LFR-based augmentation structure of Hoekstra et al., in which an interconnection matrix connects a baseline block and a learning block [18, Sec. 3.1], here a fixed one [18, Sec. 3.2]: the dynamic-parallel choice below specifies every entry of the interconnection, so nothing in it is left to learn. **[TODO: Candidate reason for the fixed interconnection (own reasoning, no source or decision entry yet); confirm or replace.]** Since Y is a state, the self-scheduled LFR of Section II-B is a nonlinear map of the state and the input, and it enters this structure as the baseline block. The augmented model is thus an LFR whose baseline block is itself an LFR in Y . The learning block is a neural network, as in Hoekstra et al., not an LPV-LFR with a learned scheduling map as in the augmentation of Drenth [19, Eq. (5.2)]. **[TODO: Missing: reason for a network as learning block rather than Drenth's LPV-LFR augmentation. The research plan (Aspect 2) extends the structure "to the LPV setting following" Drenth and leaves "the specific parameterization to be established". Candidates: (a) Hoekstra's framework with its zero start and well-posedness result is kept as is; (b) the research question asks only for a structure that captures the omitted dynamics. Close D-017 with a decision entry.]** Of the interconnections, we choose the dynamic-parallel one [18, Table 1], because the orthogonal-by-construction parametrisation that Section IV extends is defined for this additive structure [3]. It also keeps the baseline an explicit term of the state update. The output map is not augmented, because it is the exact kinematic relation (1). Kessels likewise leaves an accurate output map unaugmented [4, Remark 5.3].

The structure is a discrete-time state update, in which the learned correction is added to the baseline transition [18, Sec. 2, Table 1]. The baseline therefore enters through its transition over one sample period T_s , one fourth-order Runge-Kutta

(RK4) step of (2) with the input held over the step:

$$\begin{aligned} F(\tilde{x}, u_{\text{act}}) &= \left[M(Y)^{-1} (P u_{\text{act}} - C \dot{q} - K q) \right], \\ k_1 &= F(\tilde{x}_k, u_{\text{act},k}), \quad k_2 = F(\tilde{x}_k + \frac{T_s}{2} k_1, u_{\text{act},k}), \\ k_3 &= F(\tilde{x}_k + \frac{T_s}{2} k_2, u_{\text{act},k}), \quad k_4 = F(\tilde{x}_k + T_s k_3, u_{\text{act},k}), \\ f_{\text{base}}(\theta_{\text{base}}, \tilde{x}_k, u_{\text{act},k}) &= \tilde{x}_k + \frac{T_s}{6} (k_1 + 2k_2 + 2k_3 + k_4), \end{aligned} \quad (13)$$

where M , C and K depend on θ_{base} through (11), and $M(Y)^{-1}$ is the closed loop of the LFR of Section II-B. Each stage evaluates $M(Y)$ at the Y of its own state, so f_{base} remains self-scheduled. In physical units, the augmented model is

$$\begin{aligned} \tilde{x}_{k+1} &= f_{\text{base}}(\theta_{\text{base}}, \tilde{x}_k, u_{\text{act},k}) + f_{\text{aug}}(\theta_{\text{aug}}, \hat{x}_k, u_{\text{act},k}), \\ \bar{x}_{k+1} &= g_{\text{aug}}(\theta_{\text{aug}}, \hat{x}_k, u_{\text{act},k}), \\ \hat{y}_k &= P^\top q_k, \end{aligned} \quad (14)$$

where $\tilde{x}_k = \text{col}(q_k, \dot{q}_k) \in \mathbb{R}^6$ is the physical state of (10), $\bar{x}_k \in \mathbb{R}^{n_{\bar{x}}}$ the added state, $\hat{x}_k = \text{col}(\tilde{x}_k, \bar{x}_k)$, $u_{\text{act},k} \in \mathbb{R}^3$ the actuator forces, $\hat{y}_k \in \mathbb{R}^3$ the predicted actuator positions, whose recorded counterpart $q_{\text{act},k}$ we write y_k , $\theta_{\text{base}} \in \mathbb{R}^{10}$ the combinations (11), and $\theta_{\text{aug}} \in \mathbb{R}^{n_{\text{aug}}}$ the network parameters. The baseline transition reads only the physical state, whereas f_{aug} and g_{aug} read the full state $\hat{x}_k$: the learned part sees the added states and, through $\tilde{x}_k$, the payload position Y . No learned function feeds the baseline block or itself within a step, so the interconnection graph is acyclic [18, Thm. 1] and, with the well-posedness of the baseline LFR (Section II-B), the augmented model is well posed for every Y .

The network and the identification act on centred signals of unit variance, so that network inputs in metres, radians and newtons, and output channels of different amplitude, enter on a comparable scale. **[TODO: Candidate reason for the normalisation (own reasoning; D-055 and D-119 cover the mechanics, not the reason); confirm or replace with a source.]** Every statistic is taken from the training records, the state statistics from the positions $q = P^{-\top} y$ and their first differences divided by T_s :

$$\tilde{x}^n = T_x(\tilde{x} - \mu_x), \quad u^n = T_u(u_{\text{act}} - \mu_u), \quad y^n = T_y(y - \mu_y), \quad (15)$$

where $\mu_x \in \mathbb{R}^6$, $\mu_u, \mu_y \in \mathbb{R}^3$ are the channel means, $T_\bullet = \text{diag}(\sigma_\bullet)^{-1}$ with $\sigma_\bullet$ the channel standard deviations, and $\bar{x}$ is not scaled, so $\hat{x}^n = \text{col}(\tilde{x}^n, \bar{x})$. In these coordinates the baseline transition is $f_{\text{base}}^n(\theta_{\text{base}}, \tilde{x}^n, u^n) = T_x(f_{\text{base}}(\theta_{\text{base}}, T_x^{-1} \tilde{x}^n + \mu_x, T_u^{-1} u^n + \mu_u) - \mu_x)$: it maps back to physical units, takes the RK4 step and normalises the result. A correction only scales, $f_{\text{aug}}^n = T_x f_{\text{aug}}$, because it is a difference of states, and $g_{\text{aug}}^n = g_{\text{aug}}$. Since the position part of μ_x is $P^{-\top} \mu_y$, the output map stays linear, $\hat{y}_k^n = C^n \tilde{x}_k^n$ with $C^n = T_y P^\top [I_3 \ 0] T_x^{-1}$.

One multilayer perceptron ϕ_{aug} with tanh hidden layers and a linear output layer supplies both learned functions, $\text{col}(f_{\text{aug}}^n, g_{\text{aug}}^n) = \phi_{\text{aug}}(\theta_{\text{aug}}, \hat{x}_k^n, u_k^n)$: it reads the $6 + n_{\bar{x}}$ states and the three inputs, its first six outputs are the corrections of the physical rows and its last $n_{\bar{x}}$ outputs the next added state




Fig. 3: Augmented model. The physics step f_{base} and the network ϕ_{aug} act in parallel on $\hat{x}_k$. ϕ_{aug} supplies f_{aug} and g_{aug} , and the output map is not learned. The encoder ψ sets $\hat{x}$ at each window start. During training, the input is the closed-loop input of Section III-C. **[TODO: No setpoint generator currently for the inputs; layout cramped at one column width.]**

(Fig. 3). A feedforward network suffices because the baseline is augmented additively, whereas a series structure needs a residual network to reproduce the baseline at initialisation [18, Sec. 4.3]. The corrections act on every physical row of the baseline transition, as in the dynamic-parallel structure [18, Table 1]. The learned component thus also acts on the free axes X and Y , which have no stiffness in (4), so a persistent correction on these rows accumulates in position. **[TODO: Missing: justification for correcting the position rows rather than keeping them as the integrals of the velocities. Context: Hoekstra et al. detach the free integrators of their vehicle and learn only the input-to-velocity map, which needs measured velocities [2, Secs. 7.2, 7.4]; here only positions are measured. D-103 requires only that X and Y are reachable.]** With $n_{\bar{x}} = 0$, ϕ_{aug} corrects the physical rows only, which is the static parallel augmentation of [18, Table 1].

Because f_{aug} acts on the same rows as f_{base} , it can also learn changes that a different θ_{base} would produce, so the split between baseline and network is not unique [2, Sec. 5.2]. In the experiments of Hoekstra et al., approximately initialised physical parameters stay close to their initial values, while the network learns dynamics the baseline could represent [2, Sec. 6.3]. With a non-unique θ_{base} , the baseline can lose its physical meaning [3, Sec. 3]. Here, the estimated combinations are interpretable if they keep that meaning, which requires a unique estimate. Section IV extends the method so that this estimate is unique.

B. Encoder and initial state

- Each window starts from a state of which only the positions are measured; an encoder estimates it from

the input-output history (Beintema); the window includes sample τ (encoder paper Eqs. (11) to (17))

- The encoder is a linear map plus a nonlinear correction (encoder paper Eq. (8)); the linear part acts on scaled, uncentred histories (D-055, code)
- W^b : linearise at rest at $Y = 0$, the design point of the controller; frozen at $Y = 0$ the baseline is LTI, so the linearisation is exact there (this work)
- W^b is the reconstructability map of the encoder paper, Eqs. (15) to (17); the noise-model version (Eqs. (18) to (29)) needs a process-noise gain, our noise is measurement noise
- Adaptations: marginally stable but observable baseline; co-estimated θ_{base} (footnote 1)
- Model-based rather than data-based (2026 Eq. (30)): the baseline simulation the latter needs drifts on X and Y ; negligible cost (encoder paper Sec. 4.4) (this work)

Only the positions of the state at the start τ of a training window are measured. Following Beintema et al., an encoder estimates this state from the input and output history [20]. Simulating short windows from encoded states keeps the cost of a gradient step independent of the record length and stabilises the optimisation [20, Secs. 2, 3]. Unlike the criterion of Hoekstra et al., whose encoder reads the samples up to $\tau - 1$ [2, Eq. (22e)], the history includes sample τ , as the reconstructability map used below requires [21, Eqs. (11) to (17)]:

$$\hat{x}_{\tau|\tau}^n = \psi(\theta_{\text{enc}}, \chi_{\tau}), \quad \chi_{\tau} = \text{col}(y_{\tau-n}^n, \dots, y_{\tau}^n, u_{\tau-n}^n, \dots, u_{\tau}^n), \quad (16)$$

where n is the encoder history, $\chi_{\tau} \in \mathbb{R}^{6(n+1)}$ the recorded window and $\theta_{\text{enc}} \in \mathbb{R}^{n_{\text{enc}}}$ the encoder parameters. The encoder thus estimates all $6 + n_{\bar{x}}$ states, the measured positions included. **[TODO: Align the figure's past-only encoder-history labels with (16), which also uses the current sample.]**

The encoder is a linear map with a nonlinear correction [21, Eq. (8)]. Its added-state rows W^a have no baseline counterpart. The linear part acts on the scaled but uncentred window, for which the map (19) below is built:

$$\psi(\theta_{\text{enc}}, \chi_{\tau}) = \begin{bmatrix} W^b \\ W^a \end{bmatrix} (\chi_{\tau} + \chi_0) - \begin{bmatrix} T_x \mu_x \\ 0 \end{bmatrix} + \tilde{\psi}(\chi_{\tau}), \quad (17)$$

where $W^b \in \mathbb{R}^{6 \times 6(n+1)}$, $W^a \in \mathbb{R}^{n_{\bar{x}} \times 6(n+1)}$, χ_0 restores the means that (15) removes, $\tilde{\psi}$ is a multilayer perceptron, and θ_{enc} collects W^b , W^a and the parameters of $\tilde{\psi}$. We write ψ_a for the last $n_{\bar{x}}$ rows of ψ and θ_{enc}^a for W^a with the added-state rows of the output layer of ψ . **[TODO: VERIFY bib: hoekstra2026encoder = arXiv:2602.13108v1 (13 Feb 2026), Hoekstra, Györök, Tóth, Schoukens; matches PDF p. 1. Check for a published (IFAC) version before submission.]**

The physical rows start from the baseline, through the model-based initialisation of Hoekstra et al. [21, Sec. 3.1]. Its map needs a linear model, so as in [21, Eqs. (30) to (32)] we linearise the baseline at an equilibrium. We take rest at $Y = 0$, which is an equilibrium for zero input because K in (4) vanishes

on X and Y , and the position at which the controller of the records is designed (Section V). Frozen at $Y = 0$, (2) is linear, so the linearisation is the baseline itself at that position and approximates it only as Y moves. Discretising by zero-order hold over T_s and scaling as (15) without the means gives

$$A_d = T_x e^{A_c T_s} T_x^{-1}, \quad B_d = T_x \int_0^{T_s} e^{A_c s} ds B_c T_u^{-1}, \quad C_d = T_y C_c T_x^{-1} \\ A_c = \begin{bmatrix} 0 & I_3 \\ -M_0^{-1}K & -M_0^{-1}C \end{bmatrix}, \quad B_c = \begin{bmatrix} 0 \\ M_0^{-1}P \end{bmatrix}, \quad C_c = [P^\top \quad 0]$$
(18)

with $M_0 = M(0)$ from (5), evaluated at the nominal parameters. Stacking the outputs and inputs of the window as Y_τ and U_τ , the output and input parts of $\chi_\tau + \chi_0$ with the oldest sample first, gives $Y_\tau = O_n \tilde{x}_{\tau-n} + \Gamma_n U_\tau$ and $\tilde{x}_\tau = A_d^n \tilde{x}_{\tau-n} + r_n U_\tau$. Eliminating $\tilde{x}_{\tau-n}$ gives the reconstructability map [21, Eqs. (15) to (17)]

$$\tilde{x}_\tau = A_d^n O_n^+ (Y_\tau - \Gamma_n U_\tau) + r_n U_\tau, \\ W^b = [A_d^n O_n^+ \quad r_n - A_d^n O_n^+ \Gamma_n], \quad (19)$$

where $O_n = \text{col}(C_d, C_d A_d, \dots, C_d A_d^n) \in \mathbb{R}^{3(n+1) \times 6}$, Γ_n is block lower triangular with block (i, j) equal to $C_d A_d^{i-j-1} B_d$ for $i > j$ and zero otherwise, $r_n = [A_d^{n-1} B_d \dots B_d \quad 0]$, and $^+$ is the Moore-Penrose pseudoinverse. The paper assumes a stable baseline [21, Sec. 2]. Ours is marginally stable, but the measured positions make (A_d, C_d) observable, so O_n has full column rank. The map is derived without noise; with white Gaussian measurement noise it is the maximum-likelihood estimate of the state [21, Sec. 3.1]. The paper also gives a version for process noise, which needs a known noise gain [21, Eqs. (18) to (29)]. The noise in our records enters at the position measurement (Section V) and no noise gain is known, so we use (19). Our measurement noise is coloured, so the map is a starting value for W^b , not an optimal estimator. It leaves co-estimation of the baseline parameters out of scope [21, footnote 1], whereas here W^b is trained jointly with them. **[TODO: Fix before the final run (FINAL-RUN-FIXES.md item 1): W^b is built from the nominal parameters, which equal those of the simulated truth, while θ_{base} starts detuned. The final runs build W^b at the start $\theta_{\text{base},0}$; the text then states that.]**

Hoekstra et al. instead fit the baseline encoder to states simulated with the baseline [2, Eq. (30)], the data-based alternative of [21, Sec. 3.2]. On the gantry, that simulation is driven by the recorded input and drifts on the free axes X and Y (Section III-C), so it would fit W^b to wrong states. The map (19) needs no simulation and negligible computation [21, Sec. 4.4]. It gives better starting values and faster convergence than a random encoder, at comparable final accuracy [21, Sec. 4.4].

C. Closed-loop simulation

- The records are measured under feedback and the model must predict the machine under feedback, so each window is simulated in the loop, not driven by the recorded input (2026 Eq. (22))

- Kessels (Eqs. (5.12), (5.13); Remark 5.6, noisy records); for linear systems the indirect approach (Forssell and Ljung); free X, Y violate the stability condition behind SUBNET consistency (Beintema Cond. 1), the loop can satisfy it; motivation, not a guarantee
- Both loops displayed; subtracting the machine's controller equations gives our residual form, exact under a compatible controller, signals and initial state (check in Setup)
- One normalised closed-loop step displayed; the model's output error re-enters its input through the controller
- $x_\tau^c = 0$ needs no controller initialisation and matches the encoder; Kessels' direct form needs the machine's controller state (Remark 5.4)
- No plant feedthrough fixes the step order: no algebraic loop, no added delay; error before τ is not carried into the window

The records are measured under feedback, and the model must predict the machine under feedback: its response to a reference, a feedforward and a known controller. We therefore simulate each training window in closed loop with that controller, rather than driving the model with the recorded input as in [2, Eq. (22)]. Kessels trains augmented models in this way [4, Eqs. (5.12), (5.13)] and reports that identifying the open-loop model from his closed-loop data failed, which he attributes to noise-induced bias [4, Remark 5.6]; our records are noisy as well (Section V). For linear systems, fitting the closed loop from the reference with a known controller is the indirect approach to closed-loop identification, in which an output-error criterion remains consistent because the reference is uncorrelated with the noise [22]. The free axes give a second reason. X and Y have no stiffness, so the open-loop gantry is marginally stable: a difference in initial velocity leaves a permanent position offset. The consistency argument of the identification method [2, Sec. 5.5] rests on that of SUBNET, which assumes an incrementally exponentially output stable data-generating system [20, Cond. 1]; the open-loop gantry violates this condition, whereas the loop with the stabilising controller can satisfy it. For our nonlinear model, both arguments motivate the closed loop but do not prove consistency. **[TODO: Missing: section number for [22]. The indirect approach was verified in Sec. 5.2.1 of the LiU report LiTH-ISY-R-2021; the Automatica numbering is unchecked. Candidate: cite without a section number until the journal version is read.] [TODO: Missing: evidence on the thesis data that open-loop rollouts accumulate position error on X and Y . Candidate: the open-loop versus closed-loop drift on one record that THESIS-RESULTS.md Sec. 5 plans, reported in Section VI.]**

The controller of each record is known, linear and time invariant, and maps the position error in actuator coordinates to actuator forces. The machine and the model each close this controller around their own output, with controller states x^{fb} and $\hat{x}^{\text{fb}}$, and both loops are driven by the reference r and

feedforward u^{ff} of the record (Fig. 4):

$$\begin{aligned} x_{k+1}^{\text{fb}} &= A_{\text{fb}}x_k^{\text{fb}} + B_{\text{fb}}(r_k - y_k), \\ u_{\text{act},k} &= C_{\text{fb}}x_k^{\text{fb}} + D_{\text{fb}}(r_k - y_k) + u_k^{\text{ff}}, \\ \hat{x}_{k+1}^{\text{fb}} &= A_{\text{fb}}\hat{x}_k^{\text{fb}} + B_{\text{fb}}(r_k - \hat{y}_k), \\ \hat{u}_{\text{act},k} &= C_{\text{fb}}\hat{x}_k^{\text{fb}} + D_{\text{fb}}(r_k - \hat{y}_k) + u_k^{\text{ff}}, \end{aligned} \quad (20)$$

where $(A_{\text{fb}}, B_{\text{fb}}, C_{\text{fb}}, D_{\text{fb}})$ are the state-space matrices of the controller discretised over T_s , r_k and u_k^{ff} are in actuator coordinates, and $\hat{u}_{\text{act},k}$ is the input the model receives in place of the recorded $u_{\text{act},k}$. Subtracting the machine's lines from the model's, with $x^c = \hat{x}^{\text{fb}} - x^{\text{fb}}$, removes r and u^{ff} :

$$\begin{aligned} x_{k+1}^c &= A_{\text{fb}}x_k^c + B_{\text{fb}}(y_k - \hat{y}_k), \quad x_\tau^c = 0, \\ \hat{u}_{\text{act},k} &= u_{\text{act},k} + C_{\text{fb}}x_k^c + D_{\text{fb}}(y_k - \hat{y}_k). \end{aligned} \quad (21)$$

The model input is thus the recorded input plus the controller's response to the output error of the model. This form equals the model loop of (20) when the recorded input was produced by this controller from the same output samples and initial state. The difference of the controller states needs no estimate: setting it to zero at the window start places the model's controller in the machine's controller state, as the encoder places the model state on the recorded trajectory. Kessels' direct form instead needs the machine's controller state at every window start, read from the machine or reconstructed from the reference, the output and the controller [4, Remark 5.4]. **[TODO: Check for Setup: the training loop uses the controller re-discretised at the model rate, while the records were generated at a higher rate. Run the re-discretised controller on the recorded $r - y$ and compare its feedback force with the recorded one. A rollout with $\hat{y} = y$ gives $\hat{u} = u$ by construction and checks nothing. Related: FINAL-RUN-FIXES.md item 3, the D-232 leak check on the lowpass data, passed (0.22 to 0.50 nm).]**

The augmented model has no feedthrough, so each sample is evaluated in a fixed order: the output from the physical state, then the input, then the model and controller states advance. The controller feedthrough D_{fb} therefore closes no algebraic loop, and the loop adds no delay. In the coordinates of (15), where the output means cancel in the error, one step of a window started at τ is

$$\begin{aligned} \hat{y}_{k|\tau}^n &= C^n \hat{x}_{k|\tau}^n, \quad \varepsilon_{k|\tau} = T_y^{-1}(y_k^n - \hat{y}_{k|\tau}^n), \\ \hat{u}_{k|\tau}^n &= u_k^n + T_u(C_{\text{fb}}x_k^c + D_{\text{fb}}\varepsilon_{k|\tau}), \\ x_{k+1}^c &= A_{\text{fb}}x_k^c + B_{\text{fb}}\varepsilon_{k|\tau}, \quad x_\tau^c = 0, \end{aligned} \quad (22)$$

after which $\hat{x}_{k+1|\tau}^n$ follows from (14) with input $\hat{u}_{k|\tau}^n$, starting from the encoder state (16). Here $\varepsilon_{k|\tau} = y_k - \hat{y}_{k|\tau}$ is the output error in metres. The output error of the model thus re-enters its own input through the controller, so a model error is partly corrected within the window, as on the machine. Model error made before τ is not carried into the window, whereas the free-run validation of Section III-D applies (22) over whole records and carries it.

D. Identification problem

- θ_{base} is estimated jointly with network and encoder, as in 2026 Eq. (22); Kessels fixes the first-principles parameters (Eq. (5.12))




Fig. 4: Recorded plant loop (top) and augmented-model loop (bottom), driven by one shared reference r_k and feedforward u_k^{ff} , with the same feedback controller. The model output is evaluated before the input advances its state to the next sample.

- Training coordinates ξ , zero at the start; nine positive combinations logarithmic, the signed m_Δ through (23), so every iterate satisfies (12) without clamp or penalty (D-229)
- One display: cost over $\vartheta = (\xi, \theta_{\text{aug}}, \theta_{\text{enc}})$ with the model, encoder and loop constraints; every window sample scored; controller fixed; a closed-loop fit can hide plant error (Section VI)
- Adam with exact backpropagation through the window, then an unbounded L-BFGS polish (ξ admissible for every value); best validation free-run checkpoint; polish kept only if it improves validation and no meter regresses
- Initialisation display: $\xi = 0$, zero output layer, encoder start; reproduces the baseline (2026 Eq. (29)); $\mathcal{M}_{\text{U}}$ named

We estimate θ_{base} jointly with the network and the encoder, as Hoekstra et al. do [2, Eq. (22)], whereas Kessels leaves the first-principles parameters fixed [4, Eq. (5.12)]. Every iterate must satisfy (12), so that $M(Y) \succ 0$ and the LFR stays well posed without a clamp or a penalty. We therefore train in coordinates $\xi \in \mathbb{R}^{10}$ that are zero at the start $\theta_{\text{base},0}$. The nine positive combinations are logarithmic, $\theta_{\text{base},j} = \theta_{\text{base},0,j} e^{\xi_j}$, which keeps them positive. The signed m_Δ is bounded by (12) itself:

$$m_\Delta = \rho \frac{2}{L_b} \sqrt{m_\Sigma J_{\text{eff}}} \tanh(\eta_0 + s \xi_\Delta), \quad (23)$$

where ξ_Δ is the component of ξ for m_Δ , $\rho \in (0, 1)$ a margin, and η_0 and s reproduce $m_{\Delta,0}$ and its sensitivity at $\xi = 0$ (Appendix A-C). Every ξ thus gives $M(Y) \succ 0$ for all Y , which guarantees positive definite inertia but not positive raw masses.

Over the window starts $\mathcal{T}$ and the window length n_f , the trained parameters $\vartheta = (\xi, \theta_{\text{aug}}, \theta_{\text{enc}})$ minimise the normalised closed-loop output error, with the controller fixed. As in [2, Eq. (22)], every sample of a window is scored. We adapt that criterion in two respects: the model is simulated in the loop of Section III-C, and the encoder reads sample τ . Because the controller acts on the output error, a good closed-loop fit can mask plant error, so Section VI also tests the model under a

changed controller. The identification problem is

$$\begin{aligned} \min_{\vartheta} V(\vartheta) &= \frac{1}{|\mathcal{T}| n_f n_y} \sum_{\tau \in \mathcal{T}} \sum_{k=\tau}^{\tau+n_f-1} \|y_k^n - \hat{y}_{k|\tau}^n\|_2^2 \quad \text{s.t.} \\ \hat{x}_{k+1|\tau}^n &= f_{\text{base}}^n(\theta_{\text{base}}(\xi), \hat{x}_{k|\tau}^n, \hat{u}_{k|\tau}^n) + f_{\text{aug}}^n(\theta_{\text{aug}}, \hat{x}_{k|\tau}^n, \hat{u}_{k|\tau}^n), \\ \bar{x}_{k+1|\tau} &= g_{\text{aug}}^n(\theta_{\text{aug}}, \hat{x}_{k|\tau}^n, \hat{u}_{k|\tau}^n), \\ \hat{x}_{\tau|\tau}^n &= \psi(\theta_{\text{enc}}, \chi_{\tau}), \quad \hat{y}_{k|\tau}^n, \hat{u}_{k|\tau}^n \text{ from (22)}, \end{aligned} \quad (24)$$

for $\tau \leq k < \tau + n_f$, where $n_y = 3$ and $\theta_{\text{base}}(\xi)$ is the coordinate map above. The controller matrices are fixed, but they are part of the computation: since $\hat{u}_{k|\tau}^n$ depends on $\hat{y}_{k|\tau}^n$, the gradient of V also passes through the controller recursion.

We minimise V with Adam on minibatches of windows, with the gradient obtained by automatic differentiation through the whole window, so no part of the window is truncated from the gradient. We then polish with L-BFGS on a fixed set of windows. Drenth et al. also follow Adam with L-BFGS, in its bounded form L-BFGS-B [14, Sec. 5]. Our polish needs no bounds, because (23) and the logarithmic coordinates give an admissible baseline for every ξ . During Adam, the parameters with the lowest validation error are kept, the closed-loop free-run error over whole validation records of Section V, each simulated from one encoder start with $x^c = 0$. The polish is kept only if it lowers this error and no monitored quantity regresses. The encoder history n , the window length n_f , the added-state dimension $n_{\bar{x}}$, the width and depth of ϕ_{aug} and the optimiser settings are given with their basis in Section V. **[TODO: Fix before the final run (FINAL-RUN-FIXES.md item 2): the polish acceptance meters include `combo_err`, the error of the ten combinations against the true parameters, and possibly `param_loss`, the drift from the start. The final runs accept on the validation error and non-oracle meters only.]**

Training starts from the baseline, as Hoekstra et al. require of the initialisation [2, Eq. (29)]. Every trained block starts as

$$\begin{aligned} \xi = 0 &\implies \theta_{\text{base}} = \theta_{\text{base},0}, \\ W_L = 0, b_L = 0 &\implies f_{\text{aug}}^n = 0, g_{\text{aug}}^n = 0, \\ W^b \text{ from (19), } W^a &\sim \mathcal{U}(-a, a), \quad \tilde{W}_L = 0, \tilde{b}_L = 0, \end{aligned} \quad (25)$$

where $\theta_{\text{base},0}$ are the starting combinations (Section V), $W_L \in \mathbb{R}^{(6+n_{\bar{x}}) \times n_h}$ and $b_L \in \mathbb{R}^{6+n_{\bar{x}}}$ the weight and bias of the output layer of ϕ_{aug} with n_h its hidden width, $\tilde{W}_L$ and $\tilde{b}_L$ those of $\tilde{\psi}$, and $a = \sqrt{6/(n_{\bar{x}} + 6(n+1))}$ the Xavier bound, as in [2, Eq. (31)]. The hidden layers of ϕ_{aug} and $\tilde{\psi}$ keep the default random initialisation of the software. From the same physical initial state, the augmented model then reproduces the baseline. **[TODO: Attribution: the all-zero output layer follows the authors' public dynamic-parallel code (D-237) and is the thesis choice (D-240); [2, Sec. 5.4.3] instead starts the learning component as a linear map with random entries. Ask Jan which produced the reported S-DP results.]**

We call $\mathcal{M}_{\text{U}}$ the augmented model (14) identified by (24) from the start (25). Its input is the actuator force u_{act} and its output the actuator positions q_{act} .

E. Added-state initialisation

- Problem: from the start (25) the added states are zero after one step and V gives W^a and the added-state output rows zero gradient; V can fall through f_{aug} alone (PS2-METHOD Sec. 2.2; Section VI tests it)
- PS is this work's, adapted from two-stage regression (Hefny Sec. 2) and its use before backpropagation through time (Downey Sec. 4.2)
- S1: residual target from sample τ , as V scores it; decoder plus ψ_a ; trains θ_p , θ_{enc}^a only, own optimiser, after every Adam update
- Unexplained fraction on calibration records decides the S1 end and the fallback
- S2: once, one-step regression of g_{aug}^n on frozen encoder pairs with the recorded input, which equals the loop input when the encoder reproduces y_k
- S3: decoder discarded, (24) continues unchanged; deployed model, criterion and selection unchanged; added states latent; $\mathcal{M}_{\text{PS}}$ named

The start (25) leaves the added states without a role. They are zero after the first step, and the output depends on them only through the zero output layer, so the gradient of V with respect to W^a and to the added-state rows of W_L is zero at the start. Minimising V can then lower the error through f_{aug}^n alone, with static corrections or real lags, and leave the omitted dynamics unrepresented. Section VI tests whether it does. We therefore propose predictive state (PS) initialisation, which gives the added states a predictive role from data before the network reads them and then hands over to (24). It adapts the two-stage regression of predictive-state learning [23, Sec. 2] and its use as an initialisation before backpropagation through time [24, Sec. 4.2]. Unlike [23], the first stage predicts the residual of the model being trained rather than features of future observations, and the second stage regresses the nonlinear transition of the deployed model rather than a linear operator. **[TODO: Missing: the evidence for plain training failing in this way is from the previous campaign (PS2-METHOD.md Sec. 2.2, detuned runs, older data). Candidate: keep the sentence conditional, as written, until THESIS-RESULTS.md step 2 reports it.]**

Fig. 5 shows which samples enter each stage. Stage S1 trains the encoder's added-state outputs to predict the residual of the model being trained.

During the first updates of (24), one S1 update with its own Adam optimiser follows every update of V . S1 draws windows from the training records except a calibration subset of records. Its target is the current model's closed-loop residual over the $M = n + 1$ samples from the window start, and a temporary decoder reads only $\bar{x}_{\tau}$:

$$\begin{aligned} e_{\tau} &= \frac{1}{\sigma_e} \text{col}(y_k^n - \hat{y}_{k|\tau}^n)_{k=\tau}^{\tau+M-1}, \\ \min_{\theta_p, \theta_{\text{enc}}^a} &\frac{1}{|\mathcal{B}| M n_y} \sum_{\tau \in \mathcal{B}} \|d(\theta_p, \psi_a(\theta_{\text{enc}}, \chi_{\tau})) - e_{\tau}\|_2^2, \end{aligned} \quad (26)$$

where $\hat{y}_{k|\tau}^n$ is the output of the loop of Section III-C at the current ϑ , held fixed in the gradient, $\mathcal{B}$ a minibatch of window




Fig. 5: Windows of PS on one record, layout adapted from [23, Fig. 2]. S1 (top): the encoder maps the window χ_τ to $\hat{x}_\tau$, and the decoder d maps its added part $\bar{x}_\tau$ to the residuals e_τ of the next M samples (26). S2 (bottom): g_{aug} maps $\hat{x}_\tau$ and the recorded u_τ to the added state encoded from the window shifted by one sample (28). **[TODO: The figure shows the target from $\tau + 1$, as in the final runs (FINAL-RUN-FIXES.md item 5); the current runs start it at τ .]**

starts, σ_e the RMS of the residuals over $\mathcal{B}$, and d a decoder with parameters θ_p that maps the $n_{\bar{x}}$ added states to the Mn_y residual samples, with one tanh layer of the width of ϕ_{aug} and a zero output layer. Only θ_p and θ_{enc} receive a gradient, so the rest of ϑ is left to V . Sample τ is both the last sample of the encoder window and the first of the target, so S1 predicts exactly the residual sequence that V scores. **[TODO: Fix before the final run (FINAL-RUN-FIXES.md item 5): Hefny et al. keep history and future disjoint so that the history is uncorrelated with the noise in the future [23, Sec. 2]. Our records are noisy and the windows share sample τ , so the noise in y_τ enters both sides of (26). The final runs start the target at $\tau + 1$; the text and (26) then change accordingly.]**

Whether $\bar{x}_\tau$ predicts the residual is measured on a fixed set $\mathcal{C}$ of windows from the calibration records, by the fraction the decoder leaves unexplained,

$$\nu = \frac{\sum_{\tau \in \mathcal{C}} \|d(\theta_p, \psi_a(\theta_{\text{enc}}, \chi_\tau)) - e_\tau\|_2^2}{\sum_{\tau \in \mathcal{C}} \|e_\tau\|_2^2}. \quad (27)$$

S1 ends when ν stops falling, or after a maximum number of updates. If ν then exceeds a threshold, or already does at a fixed early update, the residual is not predictable from $\bar{x}_\tau$: S2 is skipped and training continues with the S1 changes kept. Otherwise, stage S2 fits the added-state transition so that it propagates these states. Once, at the end of S1, we encode window pairs that start at samples k_i and $k_i + 1$ anywhere in the fit records, freeze them as the set $\mathcal{P}$, and fit g_{aug} by one-step regression with a fresh optimiser, keeping the best parameters on pairs from the calibration records:

$$\min_{\theta_{\text{aug}}} \frac{\sum_{i \in \mathcal{P}} \|g_{\text{aug}}^n(\theta_{\text{aug}}, \psi(\theta_{\text{enc}}, \chi_{k_i}), u_{k_i}^n) - \psi_a(\theta_{\text{enc}}, \chi_{k_i+1})\|_2^2}{\sum_{j=1}^{n_{\bar{x}}} \sum_{i \in \mathcal{P}} (\psi_{a,j}(\theta_{\text{enc}}, \chi_{k_i+1}) - \bar{\psi}_{a,j})^2}, \quad (28)$$

minimised on minibatches of $\mathcal{P}$, where θ_{enc} is held at its value at the end of S1 and $\bar{\psi}_{a,j}$ is the mean of the j th target. The input $u_{k_i}^n$ is the recorded one: the force that moved the machine from k_i to $k_i + 1$, from whose response the target is encoded.

At a window start of the loop (21), $x^c = 0$, so the model input there is

$$\hat{u}_{\text{act},k_i} = u_{\text{act},k_i} + D_{\text{fb}}(y_{k_i} - \hat{y}_{k_i}),$$

which equals the recorded input when the encoder reproduces y_{k_i} . The loss reaches ϕ_{aug} only through its added-state rows, so S2 leaves the physical rows of W_L unchanged, while the shared hidden layers move.

In stage S3 the decoder is discarded and the minimisation of (24) continues from the modified parameters, followed by the polish and selection of Section III-D. PS adds nothing to the deployed model, the criterion or the selection rule. The added states remain latent, since PS assigns them no physical coordinates. An invertible transformation of $\bar{x}$, absorbed by ϕ_{aug} , leaves the input-output behaviour unchanged. We call $\mathcal{M}_{\text{PS}}$ the model (14) identified by (24) with this opening, from the same start as $\mathcal{M}_{\text{U}}$. **[TODO: Missing: basis of the PS switch and stop rules (S1 plateau and cap, screen update and threshold, S2 plateau and cap, snapshot sizes, calibration split as the last record of each class with three or more records). Candidate: rows of the hyperparameter table in Section V, all HEURISTIC (PS2Spec; PS2-METHOD.md Sec. 5), noting that the S1 cap, not the plateau, ended S1 in every run so far.]**

IV. INTERPRETABILITY BY ORTHOGONAL CONSTRUCTION

- Goal from Section III-A: an estimate of θ_{base} that is unique under stated conditions; names IV-A to IV-C (style profile)
- Contribution share: the third contribution of Section I, in its wording (Introduction)

The split between baseline and network in the augmented model of Section III-A is not unique, and our goal is an estimate of θ_{base} that is unique under stated conditions. To this end, we construct an orthogonal correction of the physical rows (Section IV-A) and state the identification problem of the corrected model (Section IV-B). We then give the condition under which its estimate equals the true combinations, with the scope of this result (Section IV-C). This delivers the third contribution of Section I: state-level orthogonality for the identifiable physical-parameter tangent space and its true-parameter recovery condition.

A. Orthogonal correction of the physical rows

- The non-unique split of Section III-A lets the network negate the baseline (Kessels Sec. 5.3.1.3); to first order a term along $\partial f_{\text{base}}^n / \partial \xi$ cancels a change of ξ (this work)
- On the gantry, each direction is an $M(Y)^{-1}$ -weighted mass, damping or stiffness force, propagated by RK4 (this work)
- The research plan's penalty trades orthogonality against fit; the construction needs no weight but a baseline linear in its parameters (Györök 2025 Eq. (13), Remark 1; 2026 Sec. 3.2, Eq. (2))
- Regressor: Jacobian of the normalised transition in ξ at ξ_{lin} ; range of the ten combinations; no offset column, for two reasons (Györök 2025 Sec. 4, Eqs. (15), (18); D-186; D-192; adapted)

- Fixed tuples of measured positions, differenced velocities, $\bar{x} = 0$, recorded input; data rather than simulated states (Györök 2025 Sec. 3; D-111; D-185 (?))
- Coefficient on the stack, full column rank, orthogonal corrected stack (Györök 2026 Eqs. (8), (11), Lemma 3, Assumption 1; adapted)
- Correction on the physical rows at every step; added-state rows not corrected; coefficient differentiated (Györök 2026 Eqs. (9), (13); D-190; adapted)
- Linearisation point rebuilt per epoch at a detached ξ (Györök 2025 Sec. 4, Remark 2; D-190; adapted)

The non-unique split lets the network negate the baseline, Kessels' term for learned terms that partially cancel the first-principles terms [4, Sec. 5.3.1.3]. Let $\xi' \in \mathbb{R}^{10}$ be other values of the training coordinates ξ of (23). The network term $f_{\text{aug}}^n - (\partial f_{\text{base}}^n / \partial \xi)(\xi' - \xi)$ with ξ' then gives, to first order in $\xi' - \xi$, the same one-step map as f_{aug}^n with ξ . On the gantry, each column of $\partial f_{\text{base}}^n / \partial \xi$ is the RK4 propagation of the acceleration change $-M(Y)^{-1}(\frac{\partial M}{\partial \xi_j} \ddot{q} + \frac{\partial C}{\partial \xi_j} \dot{q} + \frac{\partial K}{\partial \xi_j} q)$ of (2). A negating term therefore acts as an added mass, damping or stiffness force whose effect on the state depends on Y .

The research plan proposed the orthogonality penalty of Györök et al., whose weight trades orthogonality against fit [25, Eq. (13), Remark 1]. We instead adopt their later construction, which makes the network term orthogonal to the baseline without a weight [3, Sec. 3.2]. It needs a baseline linear in its parameters [3, Eq. (2)]. Our transition is rational in θ_{base} through $M(Y)^{-1}$, so it has no linear-in-parameters form.

We therefore linearise the normalised transition in its parameters, as Györök et al. do for a nonlinear baseline [25, Sec. 4, Eq. (15)]. We differentiate with respect to ξ rather than the fourteen raw parameters, whose four lost directions (Section II-C) would give the Jacobian four null directions. The map (23) has a nonsingular Jacobian for $m_{\Delta,0} \neq 0$, so the Jacobian in ξ spans the sensitivity directions of the ten combinations θ_{base} . Unlike their extended regressor [25, Eq. (18)], we omit the offset column f_{base}^n itself, for two reasons. The estimate of ξ moves only within the range of the Jacobian (Section IV-C), so the offset would constrain directions that do not move it. A parameter derivative also does not change when the origin of the normalised state moves, whereas the offset does, so the protected directions would depend on the centring μ_x of (15). The sensitivity is

$$\Phi(\xi_{\text{lin}}, \tilde{x}^n, u^n) = \left. \frac{\partial f_{\text{base}}^n(\theta_{\text{base}}(\xi), \tilde{x}^n, u^n)}{\partial \xi} \right|_{\xi=\xi_{\text{lin}}}, \quad (29)$$

where $\xi_{\text{lin}} \in \mathbb{R}^{10}$ is the linearisation point and $\Phi : \mathbb{R}^{10} \times \mathbb{R}^6 \times \mathbb{R}^3 \rightarrow \mathbb{R}^{6 \times 10}$. **[TODO: Missing: supervisor confirmation of the omitted offset. D-192 (the thesis OBC) gives no reason; the two reasons above are those of D-186 (2026-09-10, reasons 1 and 2), written for the earlier trajectory basis. D-190 records the supervisors' condition as non-overlap with the baseline output, which is the affine range of $[\Phi \ f_{\text{base}}^n]$. Candidate: the tangent range with these two reasons and the caveat in**

Section IV-C; parameter uniqueness needs the rank of Φ only [3, Assumption 1]. Confirm with Maarten and Jan.]

We evaluate Φ on fixed reference tuples from the training records, because only the positions of the baseline state are measured. Each tuple holds the positions $q_i = P^{-\top} q_{\text{act},i}$ of a training sample, velocities from a fourth-order central difference of q within the record, $\bar{x}_i = 0$ and the recorded input. The tuples cover every sample at which this difference exists. Each tuple is one step at the recorded input, so the controller does not enter. Györök et al. simulate the baseline when the states are not measured [25, Sec. 3]. In open loop, such a simulation drifts on the free axes (Section III-C). **[TODO: Missing: reason for data tuples rather than baseline states simulated in closed loop, which do not drift (Section III-C). D-111 rejects only open-loop simulated states, on measurement; D-185 reopened the closed-loop alternative, and no comparison is recorded. Candidate: fixed data tuples do not depend on the trained parameters, so the stack changes only when ξ_{lin} is refreshed (own reading).]**

Adapting the coefficient of Györök et al. to these tuples [3, Eqs. (8), (11), Lemma 3], we fit the network's physical rows by the stacked sensitivity in the least-squares sense. As they assume of their regressor [3, Assumption 1], the stack has full column rank:

$$\begin{aligned} \Phi(\xi_{\text{lin}}) &= \text{col}(\Phi(\xi_{\text{lin}}, \tilde{x}_i^n, u_i^n))_{i \in \mathcal{I}}, \\ f_{\text{aug}}^n(\theta_{\text{aug}}) &= \text{col}(f_{\text{aug}}^n(\theta_{\text{aug}}, \text{col}(\tilde{x}_i^n, 0), u_i^n))_{i \in \mathcal{I}}, \\ \theta_{\text{aux}}(\theta_{\text{aug}}) &= \Phi(\xi_{\text{lin}})^+ f_{\text{aug}}^n(\theta_{\text{aug}}), \\ 0 &= \Phi(\xi_{\text{lin}})^\top (f_{\text{aug}}^n(\theta_{\text{aug}}) - \Phi(\xi_{\text{lin}}) \theta_{\text{aux}}(\theta_{\text{aug}})), \end{aligned} \quad (30)$$

where $\mathcal{I}$ is the set of tuples, $N = |\mathcal{I}|$, $\tilde{x}_i^n \in \mathbb{R}^6$ and $u_i^n \in \mathbb{R}^3$ are the state and input of tuple i in the coordinates of (15), $\Phi(\xi_{\text{lin}}) \in \mathbb{R}^{6N \times 10}$, $f_{\text{aug}}^n(\theta_{\text{aug}}) \in \mathbb{R}^{6N}$, $+$ is the Moore-Penrose pseudoinverse and $\theta_{\text{aux}} \in \mathbb{R}^{10}$.

As in [3, Eqs. (9), (13)], the coefficient is a function of the network parameters, and every prediction subtracts the sensitivity times the coefficient. We subtract it from the physical rows at every step of the augmented model, at the state and input of that step. The added-state rows are not corrected, because the baseline has no added-state update to be orthogonal to. The correction is

$$f_{\text{obc}}^n(\theta_{\text{aug}}, \tilde{x}^n, u^n) = \Phi(\xi_{\text{lin}}, \tilde{x}^n, u^n) \theta_{\text{aux}}(\theta_{\text{aug}}), \quad (31)$$

where $f_{\text{obc}}^n : \mathbb{R}^{n_{\text{aug}}} \times \mathbb{R}^6 \times \mathbb{R}^3 \rightarrow \mathbb{R}^6$. Differentiating $\theta_{\text{aux}}(\theta_{\text{aug}})$ gives the network the gradient of the corrected term, which a detached coefficient would replace by that of the uncorrected one. **[TODO: Missing: the planned figure (header): the network's physical stack split into its part in range $\Phi(\xi_{\text{lin}})$ and the orthogonal part, with an inset showing the corrected physical rows, the uncorrected added-state rows and their route back through f_{aug}^n one step later. Candidate: none drawn yet.]**

We set ξ_{lin} to the current ξ at each epoch boundary, rebuild $\Phi(\xi_{\text{lin}})$ there and hold it for the epoch. No gradient therefore reaches ξ through Φ . This adapts the per-epoch refresh that Györök et al. allow in their linear case [25, Remark 2] to the linearisation point. For a nonlinear baseline, they update this

point at every evaluation of the objective or fix it at the nominal parameters [25, Sec. 4]. The stacked sensitivity changes little within an epoch, so a rebuild at every evaluation would mainly add cost. **[TODO: Missing: (a) a reported measurement for the slow change of the stacked sensitivity (D-190 measured its principal angles at 1 and 20 percent displacement; no section reports them). Candidate: one sentence with the numbers in a new appendix subsection. (b) The reason against a linearisation point fixed at the start. Candidate: the start need not be the nominal point that Györök et al. argue for, since training may start away from the true combinations (own reading).]**

B. Identification of the corrected model

- Need: the identification problem of Section III-D must include the correction (31); trained parameters, coordinates and cost stay (2026 Eq. (22); code)
- Problem display: cost and changed physical-row line, the other constraints by number; ξ_{lin} held fixed and refreshed per epoch, θ_{aux} re-solved at every evaluation (D-190 items 5, 6; code)
- Per phase: S1 fits the corrected residual; S2 does not see the correction but changes θ_{aug} ; Adam updates; validation uses the last epoch's basis; selection rebuilds the basis; the polish keeps it (D-198; code)
- $\mathcal{M}_{\text{OBC}}$ named, input and output as in Section III; compared with $\mathcal{M}_{\text{PS}}$ in Results (THESIS-RESULTS steps 5, 6; notation (?))

The identification problem of Section III-D must now include the correction in the physical-row update. The trained parameters $\vartheta = (\xi, \theta_{\text{aug}}, \theta_{\text{enc}})$ and the cost V are unchanged:

$$\begin{aligned} \min_{\vartheta} V(\vartheta) \quad \text{s.t.} \\ \tilde{x}_{k+1|\tau}^n = f_{\text{base}}^n(\theta_{\text{base}}(\xi), \tilde{x}_{k|\tau}^n, \hat{u}_{k|\tau}^n) + f_{\text{aug}}^n(\theta_{\text{aug}}, \hat{x}_{k|\tau}^n, \hat{u}_{k|\tau}^n) \\ - f_{\text{obc}}^n(\theta_{\text{aug}}, \tilde{x}_{k|\tau}^n, \hat{u}_{k|\tau}^n), \\ \text{the other constraints of (24),} \end{aligned} \quad (32)$$

where ξ_{lin} in (31) is held fixed during each update and refreshed per epoch (Section IV-A), and θ_{aux} is solved by (30) from the current θ_{aug} at every evaluation of V .

The optimiser, the stages of Section III-E and the selection rule of Section III-D are reused. Each S1 update fits the residual of the corrected model. It changes only θ_{p} and θ_{enc}^a . S2 fits g_{aug}^n , which the correction does not touch. It changes θ_{aug} , so θ_{aux} changes at the next evaluation of V . The Adam updates of V , during S1 and in S3, change every entry of ϑ . Validation uses the basis of the last epoch. After the best Adam checkpoint is restored, the basis is rebuilt at its ξ . The polish then changes every entry of ϑ with this basis, under the acceptance rule of Section III-D.

We call $\mathcal{M}_{\text{OBC}}$ the model $\mathcal{M}_{\text{PS}}$ with the correction (31), identified by (32). Its input is the actuator force u_{act} and its output the actuator positions q_{act} . It starts as $\mathcal{M}_{\text{PS}}$ does, with the first basis built at $\xi = 0$. Section VI compares $\mathcal{M}_{\text{OBC}}$ with $\mathcal{M}_{\text{PS}}$ on the ten combinations and on accuracy. **[TODO: Missing: model notation. Candidate: the calligraphic $\mathcal{M}_{\text{OBC}}$, since plain M clashes with $M(Y)$ (README open conflict on model names).]**

C. Scope and condition for true parameter recovery

- Need: (30) gives orthogonality on the tuples, not yet a unique or true estimate; that needs what the baseline misses in the data (Györök 2026 Sec. 4.1)
- True one-step discrepancy on the tuples with a successor, normalised coordinates; it also holds the reconstruction error of the tuple states (D-202, D-203; this work)
- Proposition: to first order and with full column rank, the one-step minimiser over ξ is $\xi^* + \Phi^+ \varepsilon$ for every θ_{aug} , so it does not depend on the network; true iff $\Phi^+ \varepsilon = 0$; adapted from Thm. 7 without the exact fit (this work; Györök 2026 Assumption 6, Cond. 4, Eq. (21))
- Without the correction, the minimiser moves with the network (this work)
- Scope, most consequential first: excitation condition; one-step surrogate; $\bar{x} = 0$ slice and active added states; offset caveat; normalised metric; Thms. 16, 18 do not transfer (Györök 2026 Remark 5, Secs. 4.3, 4.4; D-186, D-190, D-192, D-202)

Orthogonality on the tuples does not yet say whether the estimate is unique, nor when it equals the true combinations. Recovery depends on what the baseline misses in the data. Let $\xi^* \in \mathbb{R}^{10}$ be the coordinates of the true combinations. The successor of tuple i is tuple $i + 1$ of the same record, which every tuple except the last of each record has. On the set $\mathcal{I}' \subset \mathcal{I}$ of tuples with a successor, the true one-step discrepancy is

$$\varepsilon = \tilde{x}_+^n - f_{\text{base}}^n(\xi^*), \quad \tilde{x}_+^n = \text{col}(\tilde{x}_{i+1}^n)_{i \in \mathcal{I}'}, \quad (33)$$

where $f_{\text{base}}^n(\xi) = \text{col}(f_{\text{base}}^n(\theta_{\text{base}}(\xi), \tilde{x}_i^n, u_i^n))_{i \in \mathcal{I}'}$ and $\varepsilon, \tilde{x}_+^n \in \mathbb{R}^{6|\mathcal{I}'|}$. Since the tuple velocities are differenced positions, ε also holds the error of this reconstruction.

We adapt the recovery theorem of Györök et al., with its orthogonality condition and parameter error [3, Cond. 4, Assumption 6, Thm. 7, Eq. (21)]. Their theorem assumes that the model reproduces the data exactly. The least-squares form below needs no exact fit. It concerns the variant of (30) in which every stack runs over $\mathcal{I}'$.

Proposition 1: Suppose (i) $\Phi(\xi_{\text{lin}})$ has full column rank, and (ii) $f_{\text{base}}^n(\xi) = f_{\text{base}}^n(\xi^*) + \Phi(\xi_{\text{lin}})(\xi - \xi^*)$ for every ξ , which the gantry meets to first order in $\xi - \xi^*$ and $\xi_{\text{lin}} - \xi^*$. Then, for every θ_{aug} , the one-step error $\|\tilde{x}_+^n - f_{\text{base}}^n(\xi) - f_{\text{aug}}^n(\theta_{\text{aug}}) + \Phi(\xi_{\text{lin}})\theta_{\text{aux}}(\theta_{\text{aug}})\|_2^2$ has the same unique minimiser over ξ , $\xi = \xi^* + \Phi(\xi_{\text{lin}})^+ \varepsilon$, which does not depend on the network. It equals ξ^* if and only if $\Phi(\xi_{\text{lin}})^+ \varepsilon = 0$.

Proof: Write Φ for $\Phi(\xi_{\text{lin}})$ and $\Pi = \Phi\Phi^+$, the orthogonal projector onto range Φ . By (ii), (33) and (30), the stacked residual is $\varepsilon - \Phi(\xi - \xi^*) - (I - \Pi)f_{\text{aug}}^n$. Its projections by Π and $I - \Pi$ are $\Pi\varepsilon - \Phi(\xi - \xi^*)$ and $(I - \Pi)(\varepsilon - f_{\text{aug}}^n)$. They are orthogonal, so the one-step error is $\|\Pi\varepsilon - \Phi(\xi - \xi^*)\|_2^2 + \|(I - \Pi)(\varepsilon - f_{\text{aug}}^n)\|_2^2$. The second term does not depend on ξ . The first vanishes if and only if $\Phi(\xi - \xi^*) = \Pi\varepsilon$, which by (i) holds only for $\xi - \xi^* = \Phi^+ \varepsilon$. By (i), $\Phi^+ = (\Phi^T \Phi)^{-1} \Phi^T$, so $\Phi^+ \varepsilon = 0$ if and only if $\Phi^T \varepsilon = 0$. ■




Fig. 6: Simulated gantry with a payload-mounted absorber. The light-gray elements (m_a, k_a, c_a, δ_a) are absent from the baseline.

Without the correction, the same steps give $\xi - \xi^* = \Phi^+(\varepsilon - f_{\text{aug}}^n)$, which moves with the network along the directions of Section IV-A.

The condition $\Phi^T \varepsilon = 0$ is a condition on the excitation. Györök et al. note that an input can be designed to meet their condition when the structure of the omitted terms is known [3, Remark 5]. Without it, Proposition 1 gives a unique estimate, to first order, but not the true combinations. Because ε holds the reconstruction error, the condition can fail even where the baseline misses nothing. Proposition 1 also concerns the one-step error on the tuples, whereas (32) minimises a closed-loop multistep output error. It is therefore a surrogate, not the bias of the trained model. No orthogonality of trajectories or outputs follows from it. The tuples lie on $\bar{x} = 0$, so θ_{aux} sees the network only there. With the added states active, the part of f_{aug}^n that acts through $\bar{x}$ is not protected. The uncorrected g_{aug}^n reaches the physical rows one step later through f_{aug}^n . Without the offset column, a network term that cancels the nominal transition f_{base}^n is charged only through its component in range Φ . It does not move the estimate, because the minimiser of Proposition 1 does not depend on f_{aug}^n . The inner product is Euclidean in the coordinates of (15), so the scaling T_x weights the six rows. Finally, the consistency and covariance results of Györök et al. do not transfer, because they concern an input-output baseline linear in its parameters [3, Secs. 4.3, 4.4]. **[TODO: Missing: whether Section VI reports the overlaps $\|\Pi\varepsilon\|_2 / \|\varepsilon\|_2$ and $\|\Pi f_{\text{aug}}^n\|_2 / \|f_{\text{aug}}^n\|_2$ on held-out tuples (the old display; 06_results.tex still cites it). Candidate: not reported, since THESIS-RESULTS step 5 lists only the ten combinations and the test RMS; ε needs ξ^* , so it is available in simulation only.]**

V. EXPERIMENTAL SETUP

A. Simulation benchmark and excitation

- Generate closed-loop records from the physics baseline plus a truth-only flexible absorber, retaining exact states solely for diagnostics
- Use multisine excitation that covers the absorber pole and anti-resonance while respecting position, velocity, yaw, and force limits

- Assign independent random multisine phases to training, validation, and test records so equal spectra do not duplicate trajectories
- Include multisine-free cycloidal motion profiles matching the measured machine motion to test the operational distribution shift
- Use only the corrected all-channel Coulomb dataset matched to the frictionless excitation; the earlier mismatched-amplitude merge is superseded
- The set-valued friction law requires a velocity threshold below which a rail is treated as stuck. This is set to a physical breakaway velocity, $v_{\text{brk}} = 2.25 \cdot 10^{-3}$ m/s, taken from Lee et al. [26], Table 5, identified on a THK SSR20XW linear-motion guide. Reported Stribeck velocities for lubricated guideways span 10^{-5} to 10^{-2} m/s [26, Sec. 5.2], so the chosen value sits mid-range **[TODO: Conflicts with D-204, where the stick band is $V_{\text{EPS}} = (cc1+cc2)/m_{\text{total}} \cdot ts$ labelled THEORY: leine1998, about $3.3 \cdot 10^{-5}$ m/s before the 3 to 10x margin. Confirm which threshold the reported runs use, and whether the second citation (Stribeck range) is the same paper.]**
- State that output noise is currently disabled; add it only if a level and measurement location can be estimated from Jasper, ILC data, or stationary records
- State the sample rates, record lengths, number of windows, and total independent motion time rather than quoting overlapping-window count as data volume
- Decision (RK4 value, moved here from Section II): one reproducibility sentence giving the RK4 step, equal to the sampling period after downsampling the 20 kHz records to 4 kHz, with one substep (entry script: `fs_new=4000, up_sample=1`). Confirm against the `config.json` of each reported run, since earlier checkpoints may have used another `up_sample`
- Figure (reuse, no separate asset): let the rebuilt training-objective time strip also identify independent phases, short training windows, encoder history, and the long evaluation horizon

[TODO: Is the hidden linear absorber sufficient simulation complexity, or is a truth-only cubic spring required to demonstrate nonlinear augmentation?]

B. Real Telica data and split

- Confirm with Jasper and Dragan which reconstructed ETEL setpoint-generator or ILC-reference motions are operationally representative, including minimum/maximum move, long stroke, acceleration, and dwell
- Use the supervisor-defined split of 11 training, two validation, and two test operating points at 20 kHz
- Convert rail positions and currents to logical metres and Newtons; form velocity states by central differences because direct velocity is not logged
- Keep all iterations at one operating point in the same split to prevent identical reference trajectories leaking across sets

- Describe the X and Y operating-point grid and show whether training covers the interpolation region occupied by both test points
- Use iteration 0 for controller checks and later ILC iterations for distinct feedforward conditions without treating them as independent trajectories
- Mark real-data training and held-out evaluation as planned until a final reproducible run and force-calibration convention exist
- Figure (new): 2D operating-point grid with train, validation, and test locations, operating-range boundary, and trajectory classes; add a 3D or pairwise state-input coverage view only if it reveals coverage hidden by the 2D map

[TODO: Freeze the approved operational trajectory set with Jasper and Dragan; should direct drive velocity be requested, and does a third coverage coordinate add information beyond the 2D grid?]

C. Training protocol, comparators, and metrics

- Match records, neural seeds, windows, optimiser budget, and checkpoint criterion across baseline-only, unconstrained, tangent-OBC, affine-OBC, and black-box arms; treat multisine phase seeds separately
- Use physical RMS as the primary candidate metric for tracking and a fixed 200–400 ms post-move settling interval; use NRMS, BFR, or spectra only for a distinct supporting question
- Report physical combination error, projection overlap, added-state ablations, and encoder-init sensitivity alongside predictive metrics
- Values referenced from Section III (moved 2026-10-06), each with an honest basis (derived, inherited, tuned, compute-limited): encoder history n , window length n_f and stride, number of additional states $n_{\bar{x}}$, network width and depth, model rate
- Free-run validation (referenced from Section III-C): whole records, encoder at the record start and $x^c = 0$ once, so model error accumulates over the record unlike in a training window; state the record-length-weighted RMS used for checkpoint selection
- Decision (moved from Section II-C): joint-estimation arms start from the nominal φ of the reference implementation detuned by 10% per combination (combo_init_detune, HEURISTIC, D-191), with the nominal-parameter prior disabled (D-193); state it as a run setting with its reason
- A parameter prior bounds the drift of the combinations but does not remove the negation flat direction of Section IV; it is off in every arm (D-193)
- Practical identifiability (referenced from Section II-C): report the rank of the stacked one-step transition Jacobian with respect to φ on the training data (rank 10, nullity 4 in θ per the reduced-block docstring), from a saved artefact rather than the docstring

- Value (moved from Section II-C): the margin $\rho = 0.99$ in (23), a HEURISTIC margin inside the admissible region (D-229)
- Metric (moved from Section II-C): relative combination error in θ_{base} , with the normalisation of the small signed m_{Δ} stated explicitly
- State per arm whether raw or reduced coordinates are trained; exact-model parameter recovery is an internal implementation check, not a result (moved from Section II-C)
- Give the black-box model validation-selected capacity and horizon so it is a capability comparison rather than an artificially weakened baseline
- Evaluate held-out position/phase interpolation, ASMPt-relevant multisine-free motion, controller transfer, and operating-range extrapolation as separate rungs
- Predeclare expectation, acceptance criterion, verdict, and mechanism for every result rather than presenting isolated scores

[TODO: Decide the primary metric, credible controller-transfer arm, and feasible neural-seed count; do not pool initialization, noise, phase, and operating-point uncertainty.]

D. Reproducibility and implementation checks

- Verify controller reconstruction, coordinate transforms, normalisation, RK4 transition, and LFR equivalence before training
- Use the same controller, data conversion, and rollout entry point for training, validation, and checkpoint selection
- Controller rate (moved from Section III-C, 2026-10-06): the records are generated with the controller at 20 kHz, the training loop re-discretises it at the 4 kHz model rate (D-141); the residual form of Section III-C equals the shared-reference loop only for a compatible controller, signals and initial state
- Compatibility check for that assumption: run the re-discretised controller on the recorded $r - y$ and compare its feedback force with the recorded one; a rollout with $\hat{y} = y$ gives $\hat{u} = u$ by construction and checks nothing
- Record configuration, software precision, seed, dataset identifier, and checkpoint provenance for every reported arm
- Keep oracle states and truth parameters out of the trainable pipeline and label every diagnostic that uses them
- Figure (existing, appendix candidate): use controller-chain-v1 only as an audit trail for controller reconstruction and discretisation, not as a main method figure

[TODO: Which implementation checks belong in the main paper, and which should move to the appendix or repository documentation?]

VI. RESULTS

A. Baseline and closed-loop formulation

- Expectation and criterion: the LPV-LFR and original physics equations agree numerically across the payload

range before model-error comparisons are meaningful

- Verdict and reason: report inertia invertibility, LFR/RK4 equivalence, and the remaining baseline error caused by the truth-only absorber
- Figure: baseline and oracle output errors versus time for representative payload positions; supports the claim that the omitted mode is systematic and position dependent
- Table: baseline free-run physical RMS for LPV and frozen-LTI models across held-out Y ; tests the research plan's Aspect 1 rather than assuming LPV superiority
- Expectation and criterion: direct and residual closed-loop forms match at numerical precision when controller assumptions hold
- Evidence for Section III-C (moved 2026-10-06): on the thesis data, open-loop rollouts accumulate position error on the free axes X and Y , closed-loop rollouts do not; re-measure, do not reuse the older D-142 numbers
- Figure: recorded input, controller correction, and model error over one rollout; supports the interpretation that feedback can suppress visible output error while the open-loop model remains wrong

[TODO: The LPV-versus-frozen-LTI comparison is required by the plan but may not yet have a final matched result; run it or narrow Aspect 1 explicitly.]

B. Predictive benefit and additional-state contribution

- Expectation and criterion: augmentation must reduce held-out long-horizon error, not only the short training loss, relative to the same encoder-initialised baseline
- Verdict and reason: report the best validated improvement and the gap between window fit and free-run accuracy caused by integrated velocity bias
- Figure: validation RMS against optimiser updates for baseline and augmented arms, with window and 12 s metrics; supports or rejects objective alignment
- Figure: error spectra before and after augmentation with absorber pole and anti-resonance marked; tests whether the learned correction targets the missing band
- Figure: clamp each additional state and the full added-state vector to zero during rollout; quantifies what the trained latent dynamics add or harm
- Table: standard model, no-added-state ablation, zero-initialised latent path, and longer encoder history; connects structure changes to predictive benefit
- State the observed weak or harmful added-state contribution as a result if the final run does not reverse it, rather than inferring usefulness from state magnitude

[TODO: Are state Gramians useful for explaining each added state's contribution, or would intervention-based ablation be more direct and less linearisation-dependent?]

C. Parameter interpretability and orthogonality

- Evaluation on the validation-selected best checkpoints of R1/R2 (OBC, detuned/correct start) and R4/R5 (unconstrained, detuned/correct start): the following diagnostics

require no additional training runs; all reported predictive scores use the test set.

- Report normalised inner products with the baseline sensitivity space and the baseline-output overlap ρ_f (Eq. (??)), both on the $x_a = 0$ reference slice and with active added states along held-out test rollouts. Use the same normalised coordinates and stacking convention, and distinguish reference-slice orthogonality from observed overlap away from that slice.
- Report the ten physical combination errors relative to truth, their aggregate relative-error norm, and displacement from each start; compare OBC with unconstrained estimation for both starts.
- Report full augmented-model and standalone-baseline test prediction RMS using each checkpoint's fitted physical parameters, against the per-record noise floor, to assess physical-model usefulness alongside predictive accuracy.
- Interpret overlap and agreement between the two starts as empirical diagnostics of the decomposition. They do not prove uniqueness or true-parameter recovery; any uniqueness statement must retain the assumptions and scope of the derivation in Section IV-C.
- Expectation and criterion: OBC should reduce physical combination error without materially worsening held-out predictive fit relative to unconstrained joint estimation
- Verdict and reason: the completed arms reach 7.20% no projection, 7.69% tangent OBC, and 8.47% unfinished affine OBC from a common 9.49% start, so aggregate recovery does not establish OBC superiority
- Figure: validation fit versus combination error over training for all three arms; supports the claim that near-equal fit permits materially different physical parameters
- Figure: per-parameter error at matched budget; shows which combinations projection protects and which near-flat directions redistribute error
- Table: basis rank, condition, overlap, final and minimum parameter error, and validation RMS; links algorithm health to the recovery verdict
- Figure: ρ^* and predicted versus observed parameter bias by singular direction; supports the conditional result that true recovery is limited by non-orthogonal missing physics
- Explain why the 4.96% predicted floor accounts for part of the residual squared error but does not match the observed per-parameter pattern or prove convergence

[TODO: Should the incomplete affine arm be retained, rerun to a matched budget, or reported only as exploratory evidence?]

D. Generalisation and black-box comparison

- Expectation and criterion: improvement must progress from held-out position/phase interpolation to ASMP-relevant motion, controller transfer, and quantified operating-range extrapolation without retuning
- Verdict and reason: separate completed simulation evidence from pending Telica results and explain any failure by excitation coverage, initialisation, or objective mismatch

- Figure: tracking overview plus a dedicated 200–400 ms settling zoom for baseline, augmented model, and black box on an approved operational trajectory
- Table: primary RMS, fixed-window settling RMS, parameter availability, and training data for baseline, augmented, and black-box models on every completed validation rung
- Figure or table: train with one controller and evaluate the frozen models with a credible changed controller; pair this with a plant-domain check for feedback masking
- Figure: error versus multisine amplitude, stroke, acceleration, scheduling position, or state-input distance from the training set; quantify extrapolation rather than calling all held-out data extrapolation
- Report the black box’s early improvement and later free-run degradation as a negative comparator result if no stable matched run supersedes it
- Aggregate the final learned-model comparison across neural initialization seeds and show a representative, not best, time trace; report phase realizations separately

[**TODO:** Freeze which validation rungs are complete enough for a final claim; omit rather than speculate about unfinished controller, noise, or hardware experiments.]

VII. DISCUSSION

A. Answers to the four research aspects

- Aspect 1: state whether explicit Y scheduling earns measurable predictive value over frozen LTI baselines
- Aspect 2: distinguish the expressive value of dynamic augmentation from the observed trainability of its added states
- Aspect 3: conclude that OBC enforces a unique local split but true parameter recovery is conditional on the real discrepancy’s orthogonality and optimiser convergence
- Aspect 4: compare simulation interpolation, motion-profile transfer, amplitude extrapolation, real-data evidence, and the black-box result without merging them into one score
- Treat controller transfer as the industrial digital-twin use case and as a closed-loop distribution shift distinct from operating-range extrapolation
- Interpretability matters when physical parameters support controller redesign or health monitoring, but its practical value must be tied to a named downstream decision
- Figure plan: introduce no new figure; interpret the architecture, OBC, and result figures already presented instead of adding a summary graphic

[**TODO:** Which result is the primary answer if predictive accuracy and parameter recovery favour different arms?]

B. Limitations and lessons

- Closed-loop feedback improves numerical trainability but can hide plant mismatch, suppress orthogonality overlap, and reward a model that is poor in open loop
- Short windows underweight velocity DC bias on free integrators; varying encoder and rollout length is therefore a structural sensitivity test, not routine tuning

- One strengthened absorber, currently noiseless simulation, a finite neural-seed campaign, and truth-only dry friction bound claims about nonlinear, noisy, and dissipative missing dynamics
- State the final seed count and corrected-friction evidence here; do not retain development-era caveats after the final campaign changes them
- Central-difference velocities, force calibration, and the train-validation-test coverage grid bound the credibility of any real-data conclusion
- The research plan expected validated real-machine generalisation; any missing Telica result is a departure to disclose, not replace with design documentation
- Augmentation versus residual-FRF fitting remains a trade between trajectory-domain nonlinear prediction and a simpler frequency-domain account of missing dynamics
- Pointwise orthogonality is impossible for any nonzero missing force on this gantry: the parameter sensitivity of the three velocity rows has rank 3 at every reference sample (worst $\sigma_3/\sigma_1 = 5.7 \cdot 10^{-4}$, earlier absorber dataset), so $J^\top \Delta^* = 0$ can only arise from cancellation over the record (Section IV-C)
- Coulomb friction in the truth violates the recovery condition by construction: the c_{g1} sensitivity is proportional to $\dot{X}_1$ and the friction is $-c_{c1} \text{sign}(\dot{X}_1)$, so their inner product is $-c_{c1} \sum |\dot{X}_1| \neq 0$ for any moving record (D-204); on T_AF, OBC gives a unique split, not unbiased combinations

[**TODO:** Should a cubic-spring benchmark, noise study, and Monte Carlo analysis be completed, or should the thesis explicitly narrow its claims to the measured linear absorber case?]

C. Implications and future work

- Design excitation so $J^\top \Delta^* = 0$ when true parameter recovery, rather than only a unique decomposition, is required
- Separate learning rates for physical and neural parameters and compare convergence against the predicted nonzero OBC bias floor
- Replace the short-horizon objective or couple windows so integrated velocity bias is visible during training
- Repair and rerun the Coulomb benchmark before claiming that the method learns a legitimate nonzero-mean residual
- Extend evaluation beyond the multisine range and use machine-relevant trajectories to delimit safe interpolation and extrapolation

[**TODO:** Which future-work item is a prerequisite for the thesis claim and which genuinely belongs after graduation?]

VIII. CONCLUSION

A. Conclusions

- Answer whether LPV-LFR dynamic augmentation improves dual-gantry settling prediction over the physics baseline and black-box comparator

- State whether the learned additional states contributed measurable missing dynamics under the final validated configuration
- Conclude that state-level OBC prevents local tangent overlap by construction, while true parameters are recovered only when the physical discrepancy is orthogonal on the chosen excitation
- Give the final quantitative trade between prediction error and physical combination error on held-out data
- State whether the evidence covers relevant ASMPT motion, settling, controller transfer, or only a narrower simulation validation ladder
- Close with the practical recommendation for when ASMPT should use the augmented model and when the baseline or a simpler residual model is safer
- Figure plan: no conclusion figure; reserve the space for the final quantitative answer and limitations

[TODO: Rewrite every conclusion bullet after freezing the final result table; no planned or diagnostic result may appear as achieved.]

B. Discussion

- Why all states routing.
- why not use the encoder for the orthogonality each epoch instead of finite difference
- why not use the positions since we measure them $-i$ train nn and encoder in combination? not full answer
- black box not fair comparison, no init (BLA) no controller Thesis

Make a new folder for figures and code of the figures for the paper, and the final dataset that will be used.

* laten zien met parameters samenvoegen huidige methode niet correct ik moet iets van jacobian of iets laten zien * ik heb nog steeds niet netjes uitgewerkt voor LPV-LFR met hoe ik het heb omgeschreven die n m ik kan niet zo goed zeggen hoe het heet * verbeterde model met additional states laten zien als ik ze weg haal wat ze toevoegen = op nul zetten. ik weet niet of andere phase multisine validation en test genoeg laat zien hoe goed het werkt. * Duidelijk verantwoorden controller. vooral met integrator problem van de plant en dus X en Y geen stijfheid $K = 0$ daar. kleine integratie error propegeert gwn door, dit zei maarten ook denk dat dit de error is. * velocity central difference, why not get direct? * Zijn Gramians nuttig om per state te laten zien? * Multisine phase uitleggen kunnen we ook bronnen voor gebruiken. dat verschillende train validation en test verschillende multisine phase hebben. * Met alleen simulatie is huidige simulatie niet genoeg, cubic spring toevoegen?? * benefit van mijn closed-loop training method is ook dat ik geen encoder / states er voor hoeft te initializen. kessels wel referencen. figuur gebruiken. * Jacobian additional states * 2d / 3d grid wat jan zei van de data space of ik alles vang van het gebied van telica. * Montecarlo simulatie zoals jan? * For parameter recovery I had trajectories similar to the paper, currently not used for joint estimation? * How well do we extrapolate outside multisine range? what does jan do?

Verdedigen

* waarom geen basis functions * waarom niet zoals kessels: vooral andere literatuur kunnen zeggen waarom niet dat gebruikt, moet dat? * wat is het voordeel van interpreteerbaarheid? $-i$ controller maken, monitor health? * waarom niet frf of mass lijn fitten voor de missende dynamica * controller supresses / hides issues such as orthogonality overlap. Byproduct of closed-loop. * Hoe verdedig ik augmentation over residual frf fit?? * Hoeveelheid data kunnen verdedigen * Huidige nf window and encoder window, should i run for different lengths? * Orthogonality definition and that we orthogonalize additional states with respect to the baseline $:= 0$. additional states propagate, not included in our definition. * Kijken decisions.md * do i need to justify controller can i say what i got from my supervisor, how to approach these kind of questions?

APPENDIX A

SUPPORTING DERIVATIONS AND REPRODUCIBILITY

A. Explicit LPV-LFR realization

In (5), M_0 is (3) evaluated at $Y = 0$, and

$$M_1 = \begin{bmatrix} 0 & -m_h & 0 \\ -m_h & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix}, \quad M_2 = \begin{bmatrix} 0 & 0 & 0 \\ 0 & m_h & 0 \\ 0 & 0 & 0 \end{bmatrix}.$$

With a_1 and a_2 from (7) treated as independent signals, solving $M_0\ddot{q} + M_1a_1 + M_2a_2 = u - C\dot{q} - Kq$ for $\ddot{q}$ with $R = M_0^{-1}$ gives $\ddot{q} = R(u - Kq - C\dot{q} - M_1a_1 - M_2a_2)$. With $\tilde{x} = \text{col}(\dot{q}, \ddot{q})$, $z = \text{col}(\ddot{q}, a_1)$, and $w = \text{col}(a_1, a_2)$, the LTI part G of (10), partitioned as in [14, Eq. (6)], is

$$\begin{bmatrix} \tilde{x} \\ z \\ q \end{bmatrix} = \underbrace{\begin{bmatrix} A_x & B_w & B_u \\ C_z & D_{zw} & D_{zu} \\ C_y & D_{yw} & D_{yu} \end{bmatrix}}_G \begin{bmatrix} \tilde{x} \\ w \\ u \end{bmatrix}. \quad (34)$$

with blocks

$$\begin{aligned} A_x &= \begin{bmatrix} 0 & I_3 \\ -RK & -RC \end{bmatrix}, & B_w &= \begin{bmatrix} 0 & 0 \\ -RM_1 & -RM_2 \end{bmatrix}, & B_u &= \begin{bmatrix} 0 \\ R \end{bmatrix}, \\ C_z &= \begin{bmatrix} -RK & -RC \\ 0 & 0 \end{bmatrix}, & D_{zw} &= \begin{bmatrix} -RM_1 & -RM_2 \\ I_3 & 0 \end{bmatrix}, & D_{zu} &= \begin{bmatrix} R \\ 0 \end{bmatrix}, \\ C_y &= [I_3 \ 0], & D_{yw} &= 0, & D_{yu} &= 0. \end{aligned} \quad (35)$$

Here $\tilde{x}, z, w \in \mathbb{R}^6$ and $u, q \in \mathbb{R}^3$, so $G \in \mathbb{R}^{15 \times 15}$. The actuator frame enters only through $u = Pu_{\text{act}}$ and $q_{\text{act}} = P^\top q$ in (1).

Substituting $w = \Delta(Y)z$ into the z equation of (34) gives the algebraic loop equation $(I_6 - D_{zw}\Delta(Y))z = C_zx + D_{zu}u$. Written out,

$$\begin{bmatrix} I_3 + YRM_1 & YRM_2 \\ -YI_3 & I_3 \end{bmatrix} \begin{bmatrix} \ddot{q} \\ a_1 \end{bmatrix} = \begin{bmatrix} R(u - Kq - C\dot{q}) \\ 0 \end{bmatrix}.$$

The second block row gives $a_1 = Y\ddot{q}$, and the first then gives $M(Y)\ddot{q} = u - C\dot{q} - Kq$, that is (2), after multiplication by M_0 . The condition of [14, Def. 1], that $I_6 - D_{zw}\Delta(Y)$ is non-singular, is therefore the condition that $M(Y)$ is nonsingular: the Schur complement of the lower identity block gives

$$\begin{aligned} \det(I_6 - D_{zw}\Delta(Y)) &= \det(I_3 + YRM_1 + Y^2RM_2) \\ &= \det M(Y) / \det M_0. \end{aligned}$$

TABLE I: Nominal baseline parameters.

Parameter	Value	Parameter	Value
m_1	10.2 kg	m_2	10.7 kg
m_b	22.8 kg	m_h	10.1 kg
J_b	1.00 kg m ²	J_h	0.05 kg m ²
c_{g1}	14.5 N s/m	c_{g2}	20.3 N s/m
c_y	10.0 N s/m	d	0.10 m
c_{b1}	9.0 N m s/rad	c_{b2}	9.0 N m s/rad
k_{b1}	1987.5 N m/rad	k_{b2}	1987.5 N m/rad
L_b	0.725 m		

Since $M(Y)$ is nonsingular, the loop is solved in closed form, $\ddot{q} = M(Y)^{-1}(u - C\dot{q} - Kq)$ with

$$M(Y)^{-1} = \frac{N_0 + YN_1 + Y^2N_2}{d(Y)}, \quad d(Y) = \det M(Y),$$

where $N_0 + YN_1 + Y^2N_2 = \text{adj } M(Y)$ and $d(Y)$ are both quadratic in Y . The implementation evaluates both polynomials with Horner's scheme. **[TODO: Missing: print N_0, N_1, N_2 and $d(Y)$ explicitly? Candidate: no, appendix budget; the quadratic degree was checked symbolically (sympy, 2026-10-06).]**

[TODO: Remarks cut from Section II-B; keep here or drop?]

- Y^2 is realized by repeating Y , not as a second scheduling variable $p_2 = Y^2$, which would admit pairs with $p_2 \neq Y^2$; extra scheduling variables cause overbounding in controller synthesis [14]
- The scheduling map is a known selection of the state ($Y = q_3$), not a learned map as in the self-scheduled models of [14]
- The exact condition $\det M(Y) \neq 0$ replaces the sufficient condition of [14, Thm. 6], which needs a scheduling set in the unit ball and $\rho(D_{zw}) < 1$
- Drenth's Def. 1 is stated in discrete time but constrains only the algebraic loop equation, so it applies unchanged in continuous time

B. Nominal physical parameters

Table I lists the nominal values used by the baseline. They are taken from the ETEL gantry reported by García-Herreros et al. [1, Table I]. Here θ collects the fourteen table entries other than the fixed length L_b .

C. Positive definiteness of the inertia matrix

With $a = L_b m_{\Delta}/2$, elimination of the m_h block and completion of the square give

$$\begin{aligned} \frac{\det M(Y)}{m_h} &= (m_{\Sigma} + m_h)(J_{\text{eff}} + m_h Y^2) - (a - m_h Y)^2 \\ &= m_{\Sigma} m_h \left(Y + \frac{a}{m_{\Sigma}} \right)^2 \\ &\quad + (m_{\Sigma} + m_h) \left(J_{\text{eff}} - \frac{a^2}{m_{\Sigma}} \right). \end{aligned}$$

The 2×2 Schur complement has positive leading entry $m_{\Sigma} + m_h$. Sylvester's criterion therefore reduces positive definiteness to positivity of the completed-square expression for every Y . Hence, for positive m_h, m_{Σ} , and J_{eff} , $M(Y) \succ 0$ for all $Y \in \mathbb{R}$

if and only if (12) holds. For the physical parameterization, the difference between the two sides is

$$\begin{aligned} m_{\Sigma} J_{\text{eff}} - \frac{L_b^2}{4} m_{\Delta}^2 &= m_{\Sigma} (J_b + J_h) \\ &\quad + \frac{L_b^2}{4} [4m_1 m_2 + m_b (m_1 + m_2)], \end{aligned}$$

which is strictly positive for positive component masses and yaw inertias.

In (23), let $b(\theta_{\text{base}}) = \rho(2/L_b)\sqrt{m_{\Sigma} J_{\text{eff}}}$. The offset and scale are

$$r_0 = \frac{m_{\Delta,0}}{b(\theta_{\text{base},0})}, \quad \eta_0 = \text{atanh}(r_0), \quad s = \frac{m_{\Delta,0}}{b(\theta_{\text{base},0})(1 - r_0^2)}.$$

Here η_0 makes $\xi_{\Delta} = 0$ reproduce the initial combination, and s gives $\partial m_{\Delta}/\partial \xi_{\Delta}|_0 = m_{\Delta,0}$, matching the local relative-change scaling of the logarithmic coordinates. The construction requires $m_{\Delta,0} \neq 0$; a start near zero makes s small and the coordinate numerically inert. The exact inverse used for reporting is

$$\xi_{\Delta} = \frac{\text{atanh}(m_{\Delta}/b(\theta_{\text{base}})) - \eta_0}{s}.$$

D. OBC and experiment details

- Derive the stacked projector, rank-deficient pseudoinverse case, epoch refresh, and additional-state zero extension
- Derive $J^+ \Delta^*$ and list the measured singular-direction contributions to the predicted parameter bias
- List dataset identifiers, split membership, hyperparameters, seed, and exact metric definitions for every reported arm
- Figure (existing, optional): place controller-chain-v1 here only if the controller reconstruction cannot be audited adequately in prose and equations

[TODO: Can code and data tables be cited through a repository or DOI instead of spending paper space on full reproducibility metadata?]

REFERENCES

- [1] I. García-Herreros, X. Kestelyn, J. Gomand, R. Coleman, and P.-J. Barre, "Model-based decoupling control method for dual-drive gantry stages: A case study with experimental validations," *Control Engineering Practice*, vol. 21, no. 3, pp. 298–307, 2013.
- [2] J. H. Hoekstra, B. M. Györök, R. Tóth, and M. Schoukens, "Learning-based augmentation of first-principle models: A linear fractional representation-based approach," 2026, submitted to *Automatica*.
- [3] B. M. Györök, M. Schoukens, T. Péni, and R. Tóth, "Orthogonal-by-construction augmentation of physics-based input-output models," *IFAC Journal of Systems and Control*, vol. 35, p. 100376, 2026.
- [4] B. M. Kessels, "Ai-based model updating for nonlinear dynamical systems," PhD thesis, Eindhoven University of Technology, Eindhoven, The Netherlands, 2025.
- [5] International Roadmap for Devices and Systems (IRDS), "International roadmap for devices and systems: 2022 edition executive summary," IEEE, Tech. Rep., 2022.
- [6] T. Workman, L. Mirkarimi, J. Theil, G. Fountain, K. M. Bang, and B. Lee, "Die to wafer hybrid bonding and fine pitch considerations," in *2021 IEEE 71st Electronic Components and Technology Conference (ECTC)*. IEEE, 2021.
- [7] C. S. Teo, K. K. Tan, S. Y. Lim, S. Huang, and E. B. Tay, "Dynamic modeling and adaptive control of a h-type gantry stage," *Mechatronics*, vol. 17, no. 7, pp. 361–367, 2007.

- [8] P. Lambrechts, M. Boerlage, and M. Steinbuch, "Trajectory planning and feedforward design for electromechanical motion systems," *Control Engineering Practice*, vol. 13, no. 2, pp. 145–157, 2005.
- [9] J. Hwang and S. D. Noh, "Digital twin-based optimization of operational parameters for cluster tools in semiconductor manufacturing," *IEEE Access*, vol. 12, pp. 122 078–122 100, 2024.
- [10] J. Willard, X. Jia, S. Xu, M. Steinbach, and V. Kumar, "Integrating scientific knowledge with machine learning for engineering and environmental systems," *ACM Computing Surveys*, vol. 55, no. 4, pp. 1–37, Mar. 2022.
- [11] J. P. Noël and J. Schoukens, "Grey-box state-space identification of nonlinear mechanical vibrations," *International Journal of Control*, vol. 91, no. 5, pp. 1118–1139, 2018.
- [12] A. Retzler, R. Tóth, M. Schoukens, G. I. Beintema, J. Weigand, J.-P. Noël, Z. Kollár, and J. Swevers, "Learning-based augmentation of physics-based models: an industrial robot use case," *Data-Centric Engineering*, vol. 5, p. e12, 2024.
- [13] R. Tóth, *Modeling and Identification of Linear Parameter-Varying Systems*, ser. Lecture Notes in Control and Information Sciences. Berlin, Heidelberg: Springer, 2010, vol. 403.
- [14] R. Drenth, J. H. Hoekstra, M. Schoukens, and R. Tóth, "Efficient learning of affine and rational dependency LPV models with linear fractional representation," 2026.
- [15] H. Hong, A. Ovchinnikov, G. Pogudin, and C. Yap, "Global identifiability of differential models," *Communications on Pure and Applied Mathematics*, vol. 73, no. 9, pp. 1831–1879, 2020.
- [16] A. Ovchinnikov, A. Pillay, G. Pogudin, and T. Scanlon, "Computing all identifiable functions of parameters for ODE models," *Systems & Control Letters*, vol. 157, p. 105030, 2021.
- [17] M. Gautier and W. Khalil, "Direct calculation of minimum set of inertial parameters of serial robots," *IEEE Transactions on Robotics and Automation*, vol. 6, no. 3, pp. 368–373, 1990.
- [18] J. H. Hoekstra, C. Verhoek, R. Tóth, and M. Schoukens, "Learning-based model augmentation with LFRs," *European Journal of Control*, vol. 86, p. 101304, 2025.
- [19] R. Drenth, "Efficient gradient-based learning of LPV models with linear fractional representation," Master's thesis, Eindhoven University of Technology, 2025.
- [20] G. I. Beintema, M. Schoukens, and R. Tóth, "Deep subspace encoders for nonlinear system identification," *Automatica*, vol. 156, p. 111210, 2023.
- [21] J. H. Hoekstra, B. M. Györök, R. Tóth, and M. Schoukens, "Encoder initialisation methods in the model augmentation setting," 2026.
- [22] U. Forssell and L. Ljung, "Closed-loop identification revisited," *Automatica*, vol. 35, no. 7, pp. 1215–1241, 1999.
- [23] A. Hefny, C. Downey, and G. J. Gordon, "Supervised learning for dynamical system learning," in *Advances in Neural Information Processing Systems*, vol. 28, 2015.
- [24] C. Downey, A. Hefny, B. Li, B. Boots, and G. J. Gordon, "Predictive state recurrent neural networks," in *Advances in Neural Information Processing Systems*, vol. 30, 2017.
- [25] "Source pending verification," citation metadata not yet verified.
- [26] C.-Y. Lee, S.-H. Hwang, E. Nam, and B.-K. Min, "Identification of mass and sliding friction parameters of machine tool feed drive using recursive least squares method," *The International Journal of Advanced Manufacturing Technology*, vol. 109, no. 9-12, pp. 2831–2844, 2020.